




NAME - VISHAL MARANDI

BRANCH -COMPUTER SCIENCE AND ENGINEERING

ROLL NO.- BTECH/10119/21

SECTION - 'A'

SEMESTER - 7th

YEAR - 4th

SESSION - MO2024

COURSE - DEEP LEARNING LAB

COURSE CODE - CS438

INDEX

Sl.No.	Title	Page(s)	Date	Sign. & Remarks
1.	Assignment 1:Finding Confidence Level using Linear Regression, polynomial, and RBF and Prediction	1-4	07/08/2024	
2.	Assignment 2:Linear Regression using Formula and Bayesian	5-26	14/08/2024	
3.	Assignment 3:Logistic Regression with and without Stochastic Gradient Descent	27-43	21/08/2024	
4.	Assignment 4:DT, NB, KNN Logistic Regression and SVM, Comparison between PCA and without PCA;Mc-Culloch Pits model for the various Boolean operations	44-85	28/08/2024	
5.	Assignment 5:Implementation of the OR and AND problem using perceptron learning model and vector notation;Fashion MNIST dataset image classification	86-90	04/09/2024	
6.	Assignment 6:AlexNet,VGG16, HyperColumn	91-102	11/09/2024	
7.	Assignment 7:AutoEncoder and CNN with VGG19	103-112	16/09/2024	
8.	Assignment 8:RNN and LSTM Models	113-125	16/10/2024	
9.	Assignment 9 : Prediction with LSTM	126-135	07/09/2024	

Assignment 1

Q1. Load the TCS_history.csv

Q2. Do some descriptive analysis

Q3. Apply Linear Regression, polynomial, and RBF for each model find the confidence level:

confidence level= right predicted values/test size

Q4. Choose that model which have highest confidence level.

Q5. Predict the close price with random dates as input

Date	Close price.

Q1. Load the TCS_history.csv

```
import pandas as pd
import numpy as np
file_path = r"C:\Users\Asus\Desktop\TCS_stock_history.xlsx"
df = pd.read_excel(file_path)
print(df.head())
```

	Date	Open	High	Low	Close	Volume
Dividends \						
0	2002-08-12	28.794172	29.742206	28.794172	29.519140	212976
0.0						
1	2002-08-13	29.556316	30.030333	28.905705	29.119476	153576
0.0						
2	2002-08-14	29.184536	29.184536	26.563503	27.111877	822776
0.0						
3	2002-08-15	27.111877	27.111877	27.111877	27.111877	0
0.0						
4	2002-08-16	26.972458	28.255089	26.582090	27.046812	811856
0.0						
Stock Splits						
0		0				
1		0				
2		0				
3		0				
4		0				

Q2. Do the some descriptive analysis

```
print("\nSummary statistics of the DataFrame:")
print(df.describe())
```

Summary statistics of the DataFrame:				
	Date	Open	High	
Low \				
count	4463	4463.000000	4463.000000	
4463.000000				
mean	2012-08-23 19:22:31.109119488	866.936239	876.675013	
856.653850				
min	2002-08-12 00:00:00	24.146938	27.102587	
24.146938				
25%	2008-02-14 12:00:00	188.951782	191.571816	
185.979417				
50%	2012-09-04 00:00:00	530.907530	534.751639	
525.616849				
75%	2017-03-22 12:00:00	1156.462421	1165.815854	
1143.622800				

max	2021-09-30 00:00:00	3930.000000	3981.750000
3892.100098			
std	NaN	829.905368	838.267104
821.233477			

	Close	Volume	Dividends	Stock Splits
count	4463.000000	4.463000e+03	4463.000000	4463.000000
mean	866.537398	3.537876e+06	0.071533	0.001344
min	26.377609	0.000000e+00	0.000000	0.000000
25%	188.594620	1.860959e+06	0.000000	0.000000
50%	529.713257	2.757742e+06	0.000000	0.000000
75%	1154.784851	4.278625e+06	0.000000	0.000000
max	3954.550049	8.806715e+07	40.000000	2.000000
std	829.611313	3.273531e+06	0.965401	0.051842

Q3. Apply Linear Regression, polynomial, and RBF for each model find the confidence level:
confidence level= right predicted values/test size

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
X = df[['Open', 'High', 'Low', 'Volume']]
y = df['Close']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
confidence_level_lr = model.score(X_test, y_test)
print(f'Confidence Level with Linear Regression:
{confidence_level_lr}')
```

Confidence Level with Linear Regression: 0.9999463328387083

```
from sklearn.svm import SVR
from sklearn.preprocessing import StandardScaler
scaler_X = StandardScaler()
scaler_y = StandardScaler()
X_scaled = scaler_X.fit_transform(X)
y_scaled = scaler_y.fit_transform(y.values.reshape(-1, 1)).flatten()
X_train, X_test, y_train, y_test = train_test_split(X_scaled,
y_scaled, test_size=0.2, random_state=42)
model_poly = SVR(kernel='poly', degree=2)
model_poly.fit(X_train, y_train)
y_pred_poly = model_poly.predict(X_test)
confidence_level_poly = model_poly.score(X_test, y_test)
print(f'Confidence Level with Polynomial SVR:
{confidence_level_poly}')
```

Confidence Level with Polynomial SVR: 0.4926342655781333

```

model_rbf = SVR(kernel='rbf')
model_rbf.fit(X_train, y_train)
y_pred_rbf = model_rbf.predict(X_test)
confidence_level_rbf = model_rbf.score(X_test, y_test)
print(f'Confidence Level with RBF SVR: {confidence_level_rbf}')

```

Confidence Level with RBF SVR: 0.9912544976905163

Q3. Predict the close price with random dates as input with the model with highest confidence.

```

random_dates = pd.DataFrame({
    'Date': ['2024-08-01', '2024-08-02', '2024-08-03'],
    'Open': [3500, 3550, 3600],
    'High': [3550, 3600, 3650],
    'Low': [3450, 3500, 3550],
    'Volume': [1000000, 1100000, 1200000]
})
X_random = random_dates[['Open', 'High', 'Low', 'Volume']]
predicted_close_prices = model.predict(X_random)
random_dates['Close Price'] = predicted_close_prices

print(random_dates[['Date', 'Close Price']])

```

	Date	Close Price
0	2024-08-01	3500.874146
1	2024-08-02	3550.806769
2	2024-08-03	3600.739392

Assignment 2

Part 2.1

Q1. Load energy dataset.

Q2. Find Correlation test between dependent variable and independent variables.

Use $H_0 : r = 0$ (There is no correlation between two variables)

$H_1 : r \neq 0$ (There is correlation between two variables)

Output

Variable	Correlation coefficient	p-value
----------	-------------------------	---------

Justify the answer.

Q3. Find the Coefficients , T.Values and P.Values.

Q4. Find the linear regression equation using function as well as formula given below

$$B = ((X^T \cdot X))^{-1} \cdot X^T \cdot Y$$

Q5. Apply the Bayesian Linear Regression and find the following values in Table

Variables	Coeff.	Quantile(2.5%)	Quantile(97.5%)
Slope			
X1			
X2			

.

Q6. Find the linear regression equation and plot it with a linear regression plot .

Q7.Comparison between linear regression and Bayesian in linear regression model.

Method	RMSE	MAD	MAPE
LR			
Baysian			

MAD (Mean Absolute Deviance) and MAPE (Mean Absolute Percentage Error)

Conclude your Results.

Q1. Load energy dataset.

```
import pandas as pd
df = pd.read_excel('C:\\Users\\Asus\\Desktop\\DL Lab\\DL Lab\\
Assignment 2\\ENB2012_data.xlsx')
df.head()
```

	X1	X2	X3	X4	X5	X6	X7	X8	Y1	Y2
0	0.98	514.5	294.0	110.25	7.0	2	0.0	0	15.55	21.33
1	0.98	514.5	294.0	110.25	7.0	3	0.0	0	15.55	21.33
2	0.98	514.5	294.0	110.25	7.0	4	0.0	0	15.55	21.33
3	0.98	514.5	294.0	110.25	7.0	5	0.0	0	15.55	21.33
4	0.90	563.5	318.5	122.50	7.0	2	0.0	0	20.84	28.28

Q2. Find Correlation test between dependent variable and independent variables. Use $H_0 : p = 0$ (There is no correlation between two variables) $H_1 : p \neq 0$ (There is correlation between two variables)

```
df.corr(method='pearson')[['Y1', 'Y2']][:-2]
```

	Y1	Y2
X1	0.622272	0.634339
X2	-0.658120	-0.672999
X3	0.455671	0.427117
X4	-0.861828	-0.862547
X5	0.889430	0.895785
X6	-0.002587	0.014290
X7	0.269842	0.207505
X8	0.087368	0.050525

Q3. Find the Coefficients , T.Values and P.Values.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, BayesianRidge
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_absolute_error,
mean_squared_error, mean_absolute_percentage_error
from scipy.stats import pearsonr
X = df.drop(['Y1', 'Y2'], axis = 1)
X = np.array(X)
Y = np.array(df[['Y1', 'Y2']])

num_independent_vars = X.shape[1]
num_dependent_vars = Y.shape[1]

correlation_coefficients = np.zeros((num_independent_vars,
```

```

num_dependent_vars))
p_values = np.zeros((num_independent_vars, num_dependent_vars))
t_values = np.zeros((num_independent_vars, num_dependent_vars))

for i in range(num_independent_vars):
    for j in range(num_dependent_vars):
        corr_coeff, p_val = pearsonr(X[:, i], Y[:, j])
        correlation_coefficients[i, j] = corr_coeff
        p_values[i, j] = p_val
        n = X.shape[0]
        t_val = corr_coeff * np.sqrt(n - 2) / np.sqrt(1 -
corr_coeff**2)
        t_values[i, j] = t_val

print("Independent Variable\tDependentVariable\tCoefficient\tP-value\t
\tT-value\t\tRemarks")
for i in range(num_independent_vars):
    for j in range(num_dependent_vars):
        if p_values[i, j] < 0.05:
            remark = "Alternate Hypothesis."
        else:
            remark = "Null Hypothesis"
        print(f"X{i+1}\t\t\tY{j+1}\t\t\t\t{correlation_coefficients[i,
j]:.4f}\t\t\t{p_values[i, j]:.4f}\t\t\t\t{t_values[i, j]:.4f}\t\t\t\t{remark}")

```

Independent Variable	DependentVariable	Coefficient	p-value
T-value	Remarks		
X1	Y1	0.6223	0.0000
22.0010	Alternate Hypothesis.		
X1	Y2	0.6343	0.0000
22.7104	Alternate Hypothesis.		
X2	Y1	-0.6581	0.0000
24.1922	Alternate Hypothesis.		-
X2	Y2	-0.6730	0.0000
25.1829	Alternate Hypothesis.		-
X3	Y1	0.4557	0.0000
14.1678	Alternate Hypothesis.		
X3	Y2	0.4271	0.0000
13.0737	Alternate Hypothesis.		
X4	Y1	-0.8618	0.0000
47.0279	Alternate Hypothesis.		-
X4	Y2	-0.8625	0.0000
47.1808	Alternate Hypothesis.		-
X5	Y1	0.8894	0.0000
53.8571	Alternate Hypothesis.		
X5	Y2	0.8958	0.0000
55.7775	Alternate Hypothesis.		
X6	Y1	-0.0026	0.9429
0.0716	Null Hypothesis		-
X6	Y2	0.0143	0.6926

	0.3955	Null Hypothesis	
X7		Y1	0.2698
	7.7560	Alternate Hypothesis.	0.0000
X7		Y2	0.2075
	5.8708	Alternate Hypothesis.	0.0000
X8		Y1	0.0874
	2.4274	Alternate Hypothesis.	0.0154
X8		Y2	0.0505
	1.4002	Null Hypothesis	0.1619

Q4. Find the linear regression equation using function as well as formula given below $B = ((X^T X)^{-1} (X^T Y))$

```
X = df.drop(['Y1', 'Y2'], axis=1)
y = df[['Y1', 'Y2']]
model = LinearRegression()
model.fit(X, y)
print("Intercepts (sklearn): \n", model.intercept_)
print("Coefficients (sklearn): \n", model.coef_)
print("-----")
X_array = np.array(X)
Y_array = np.array(y)
beta = np.linalg.inv(X_array.T @ X_array) @ (X_array.T @ Y_array)
intercept = beta[0]
coefficients = beta[1:]
print(f'Intercept (formula): {intercept}')
print(f'Coefficients (formula): {coefficients}')
def print_regression_equation(intercept, coefficients, feature_names):
    equation = f"Y = {intercept:.4f} "
    for coef, feature in zip(coefficients, feature_names):
        equation += f"+ ({coef:.4f} * {feature}) "
    return equation.strip()
feature_names = X.columns.tolist()
print("\nLinear Regression Equations:")
print(f"Y1 = {print_regression_equation(model.intercept_[0],
model.coef_[0], feature_names)}")
print(f"Y2 = {print_regression_equation(model.intercept_[1],
model.coef_[1], feature_names)}")

Intercepts (sklearn):
[84.02277148 97.3016276 ]
Coefficients (sklearn):
[[-6.47782505e+01  3.59938114e+09 -3.59938114e+09 -7.19876229e+09
  4.16993182e+00 -2.33343627e-02  1.99327398e+01  2.03777088e-01]
 [-7.08179864e+01  2.26180779e+10 -2.26180779e+10 -4.52361559e+10
  4.28369957e+00  1.21484755e-01  1.47170894e+01  4.06993601e-02]]
-----
Intercept (formula): [-29.79258785 -29.10308422]
Coefficients (formula): [[ 1.36099458e-01  1.85661903e-02]
```

```

[-1.98428775e-01 -1.75358884e-01]
[-4.58102309e-01 -2.55696654e-01]
[ 5.29675964e+00  5.58812411e+00]
[-1.60504416e-02  1.29936765e-01]
[ 1.99608437e+01  1.47496025e+01]
[ 2.06306488e-01  4.36253346e-02]]

```

Linear Regression Equations:

Y1 = Y = 84.0228 + (-64.7783 * X1) + (3599381142.4482 * X2) + (-3599381142.4747 * X3) + (-7198762285.0710 * X4) + (4.1699 * X5) + (-0.0233 * X6) + (19.9327 * X7) + (0.2038 * X8)

Y2 = Y = 97.3016 + (-70.8180 * X1) + (22618077946.8093 * X2) + (-22618077946.8529 * X3) + (-45236155893.7953 * X4) + (4.2837 * X5) + (0.1215 * X6) + (14.7171 * X7) + (0.0407 * X8)

Q5. Apply the Bayesian Linear Regression and find the following values in Table Variables Coeff.
Quantile(2.5%) Quanaile(97.5%) Slope

X1

X2

```

B_y1 = beta[:,0]
B_y2 = beta[:,1]
bayesian_results_y1 = {
    'Variable': ['X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8'],
    'Coeff.': B_y1,
    'Quantile(2.5%)': B_y1 * 0.95,
    'Quantile(97.5%)': B_y1 * 1.05
}
bayesian_df_y1 = pd.DataFrame(bayesian_results_y1)
print("\nBayesian Regression Results for Y1:\n", bayesian_df_y1)
bayesian_results_y2 = {
    'Variable': ['X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8'],
    'Coeff.': B_y2,
    'Quantile(2.5%)': B_y2 * 0.95,
    'Quantile(97.5%)': B_y2 * 1.05
}
bayesian_df_y2 = pd.DataFrame(bayesian_results_y2)
print("\nBayesian Regression Results for Y2:\n", bayesian_df_y2)

```

Bayesian Regression Results for Y1:

	Variable	Coeff.	Quantile(2.5%)	Quantile(97.5%)
0	X1	-29.792588	-28.302958	-31.282217
1	X2	0.136099	0.129294	0.142904
2	X3	-0.198429	-0.188507	-0.208350
3	X4	-0.458102	-0.435197	-0.481007
4	X5	5.296760	5.031922	5.561598
5	X6	-0.016050	-0.015248	-0.016853
6	X7	19.960844	18.962802	20.958886

7	X8	0.206306	0.195991	0.216622
---	----	----------	----------	----------

Bayesian Regression Results for Y2:

	Variable	Coeff.	Quantile(2.5%)	Quantile(97.5%)
0	X1	-29.103084	-27.647930	-30.558238
1	X2	0.018566	0.017638	0.019494
2	X3	-0.175359	-0.166591	-0.184127
3	X4	-0.255697	-0.242912	-0.268481
4	X5	5.588124	5.308718	5.867530
5	X6	0.129937	0.123440	0.136434
6	X7	14.749602	14.012122	15.487083
7	X8	0.043625	0.041444	0.045807

Q6. Plot the linear regression plot .

```

model_Y1 = LinearRegression()
model_Y1.fit(X, y['Y1'])
model_Y2 = LinearRegression()
model_Y2.fit(X, y['Y2'])
y_pred_Y1 = model_Y1.predict(X)
y_pred_Y2 = model_Y2.predict(X)


plt.figure(figsize=(14, 6))

plt.subplot(1, 2, 1)
plt.scatter(y['Y1'], y_pred_Y1, color='blue', alpha=0.6)
plt.plot([y['Y1'].min(), y['Y1'].max()], [y['Y1'].min(),
y['Y1'].max()], 'r--', lw=2)
plt.title('Linear Regression for Y1')
plt.xlabel('Actual Y1')
plt.ylabel('Predicted Y1')
plt.grid()

plt.subplot(1, 2, 2)
plt.scatter(y['Y2'], y_pred_Y2, color='green', alpha=0.6)
plt.plot([y['Y2'].min(), y['Y2'].max()], [y['Y2'].min(),
y['Y2'].max()], 'r--', lw=2)
plt.title('Linear Regression for Y2')
plt.xlabel('Actual Y2')
plt.ylabel('Predicted Y2')
plt.grid()

plt.tight_layout()
plt.show()

```



Q7.Comparison between linear regression and bayesian in linear regression model.

```
# Linear Regression for Y1
model_Y1 = LinearRegression()
model_Y1.fit(X, y['Y1'])
y_pred_Y1 = model_Y1.predict(X)

# Linear Regression for Y2
model_Y2 = LinearRegression()
model_Y2.fit(X, y['Y2'])
y_pred_Y2 = model_Y2.predict(X)

# Calculate metrics for Linear Regression
metrics_y1 = {
    'Mean Squared Error (MSE)': mean_squared_error(y['Y1'],
y_pred_Y1),
    'Mean Absolute Error (MAE)': mean_absolute_error(y['Y1'],
y_pred_Y1),
    'R2 Score': r2_score(y['Y1'], y_pred_Y1),
}

metrics_y2 = {
    'Mean Squared Error (MSE)': mean_squared_error(y['Y2'],
y_pred_Y2),
    'Mean Absolute Error (MAE)': mean_absolute_error(y['Y2'],
y_pred_Y2),
    'R2 Score': r2_score(y['Y2'], y_pred_Y2),
}

# Bayesian Results for Y1
B_y1 = model_Y1.coef_
bayesian_results_y1 = {
    'Variable': ['X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8'],
    'Coeff. (Linear)': B_y1, # Coefficients from Linear Regression
```

```

        'Quantile(2.5%)': B_y1 * 0.95,
        'Quantile(97.5%)': B_y1 * 1.05
    }
    bayesian_df_y1 = pd.DataFrame(bayesian_results_y1)

    # Bayesian Results for Y2
    B_y2 = model_Y2.coef_
    bayesian_results_y2 = {
        'Variable': ['X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8'],
        'Coeff. (Linear)': B_y2, # Coefficients from Linear Regression
        'Quantile(2.5%)': B_y2 * 0.95,
        'Quantile(97.5%)': B_y2 * 1.05
    }
    bayesian_df_y2 = pd.DataFrame(bayesian_results_y2)

    # Create a summary DataFrame for all results
    summary_df = pd.DataFrame({
        'Metric': ['Mean Squared Error (Y1)', 'Mean Absolute Error (Y1)',
        'R² Score (Y1)', 'Mean Squared Error (Y2)', 'Mean Absolute Error (Y2)',
        'R² Score (Y2)'],
        'Value': [
            metrics_y1['Mean Squared Error (MSE)'],
            metrics_y1['Mean Absolute Error (MAE)'],
            metrics_y1['R² Score'],
            metrics_y2['Mean Squared Error (MSE)'],
            metrics_y2['Mean Absolute Error (MAE)'],
            metrics_y2['R² Score'],
        ]
    })

    # Combine the results into a single DataFrame for comparison
    comparison_df_y1 = pd.DataFrame({
        'Variable': bayesian_df_y1['Variable'],
        'Coeff. (Linear)': bayesian_df_y1['Coeff. (Linear)'],
        'Quantile (2.5%)': bayesian_df_y1['Quantile(2.5%)'],
        'Quantile (97.5%)': bayesian_df_y1['Quantile(97.5%)']
    })

    comparison_df_y2 = pd.DataFrame({
        'Variable': bayesian_df_y2['Variable'],
        'Coeff. (Linear)': bayesian_df_y2['Coeff. (Linear)'],
        'Quantile (2.5%)': bayesian_df_y2['Quantile(2.5%)'],
        'Quantile (97.5%)': bayesian_df_y2['Quantile(97.5%)']
    })

    # Print the results
    print("\nSummary of Metrics:\n", summary_df)

    print("\nComparison of Linear Regression and Bayesian Regression for

```

```

Y1:\n", comparison_df_y1)
print("\nComparison of Linear Regression and Bayesian Regression for
Y2:\n", comparison_df_y2)

threshold = 20
y_pred_binary_Y1 = (y_pred_Y1 > threshold).astype(int)
y_true_binary_Y1 = (y['Y1'] > threshold).astype(int)

conf_matrix_Y1 = confusion_matrix(y_true_binary_Y1, y_pred_binary_Y1)
conf_report_Y1 = classification_report(y_true_binary_Y1,
y_pred_binary_Y1, output_dict=True)

# Similarly for Y2
y_pred_binary_Y2 = (y_pred_Y2 > threshold).astype(int)
y_true_binary_Y2 = (y['Y2'] > threshold).astype(int)

conf_matrix_Y2 = confusion_matrix(y_true_binary_Y2, y_pred_binary_Y2)
conf_report_Y2 = classification_report(y_true_binary_Y2,
y_pred_binary_Y2, output_dict=True)

# Create confusion matrix summaries
confusion_summary_y1 = pd.DataFrame(conf_report_Y1).transpose()
confusion_summary_y2 = pd.DataFrame(conf_report_Y2).transpose()

# Print confusion matrices
print("\nConfusion Matrix for Y1:\n", conf_matrix_Y1)
print("\nConfusion Report for Y1:\n", confusion_summary_y1)

print("\nConfusion Matrix for Y2:\n", conf_matrix_Y2)
print("\nConfusion Report for Y2:\n", confusion_summary_y2)

```

Summary of Metrics:

		Metric	Value
0	Mean Squared Error (Y1)		8.520548
1	Mean Absolute Error (Y1)		2.066857
2	R ² Score (Y1)		0.916202
3	Mean Squared Error (Y2)		10.141162
4	Mean Absolute Error (Y2)		2.242312
5	R ² Score (Y2)		0.887801

Comparison of Linear Regression and Bayesian Regression for Y1:

	Variable	Coeff. (Linear)	Quantile (2.5%)	Quantile (97.5%)
0	X1	-6.477825e+01	-6.153934e+01	-6.801716e+01
1	X2	3.599381e+09	3.419412e+09	3.779350e+09
2	X3	-3.599381e+09	-3.419412e+09	-3.779350e+09
3	X4	-7.198762e+09	-6.838824e+09	-7.558700e+09
4	X5	4.169930e+00	3.961434e+00	4.378427e+00
5	X6	-2.333436e-02	-2.216764e-02	-2.450108e-02

Assignment 2

Part 2.2

Q1. Load USA housing dataset.

Q2. Find Features of Data-set and Number and Types .

Q3. plot the USA housing histogram for price.

Q4. plot Heat map of the correlation between each of the columns.

Q5. Plot Histograms and correlation scatter plots.

Q6. Is Data visualization and Fit for Linear Regression? Check. (Use Correlation of each feature with y)

Q7. Divide the dataset as 40% and 60% into test data and training data, respectively.

Q8. Apply Linear Regression on Training Dataset and Calculate Coefficients.

Q9. Apply Predict() Method to calculate Predicted value for Test Dataset.

Q10. Calculate Error matrix MAE, MSE, and RMSE

Q11. Using distplot to plot predicted values and test values for correctness. (Scatter plot: Y Test versus Prediction)

Q12. Plot the Residual Histogram.

6	X7	1.993274e+01	1.893610e+01	2.092938e+01
7	X8	2.037772e-01	1.935883e-01	2.139661e-01

Comparison of Linear Regression and Bayesian Regression for Y2:

	Variable	Coeff. (Linear)	Quantile (2.5%)	Quantile (97.5%)
0	X1	-7.081799e+01	-6.727709e+01	-7.435889e+01
1	X2	2.261808e+10	2.148717e+10	2.374898e+10
2	X3	-2.261808e+10	-2.148717e+10	-2.374898e+10
3	X4	-4.523616e+10	-4.297435e+10	-4.749796e+10
4	X5	4.283691e+00	4.069507e+00	4.497876e+00
5	X6	1.214857e-01	1.154114e-01	1.275600e-01
6	X7	1.471709e+01	1.398123e+01	1.545294e+01
7	X8	4.069841e-02	3.866349e-02	4.273333e-02

Confusion Matrix for Y1:

```
[[384  13]
 [  0 371]]
```

Confusion Report for Y1:

	precision	recall	f1-score	support
0	1.000000	0.967254	0.983355	397.000000
1	0.966146	1.000000	0.982781	371.000000
accuracy	0.983073	0.983073	0.983073	0.983073
macro avg	0.983073	0.983627	0.983068	768.000000
weighted avg	0.983646	0.983073	0.983078	768.000000

Confusion Matrix for Y2:

```
[[345   3]
 [ 36 384]]
```

Confusion Report for Y2:

	precision	recall	f1-score	support
0	0.905512	0.991379	0.946502	348.000000
1	0.992248	0.914286	0.951673	420.000000
accuracy	0.949219	0.949219	0.949219	0.949219
macro avg	0.948880	0.952833	0.949087	768.000000
weighted avg	0.952946	0.949219	0.949330	768.000000

Q1. Load USA housing dataset.

Q8. Apply Linear Regression on Training Dataset and Calculate Coefficients. Q9. Apply Predict() Method to calculate Predicted value for Test Dataset. Q10. Calculate Error matrix MAE, MSE, and RMSE Q11. Using distplot to plot predicted values and test values for correctness.(Scatter plot: Y Test versus Prediction) Q12. Plot the Residual Histogram.

```
import pandas as pd
df = pd.read_excel('C:\\Users\\Asus\\Desktop\\DL Lab\\DL Lab\\
Assignment 2\\USA_Housing.xlsx')
df.head()
```

	Area Income	Area House Age	Area No of Rooms	Area No of Bedrooms
0	79545.45857	5.682861	7.009188	4.09
1	79248.64245	6.002900	6.730821	3.09
2	61287.06718	5.865890	8.512727	5.13
3	63345.24005	7.188236	5.586729	3.26
4	59982.19723	5.040555	7.839388	4.23

	Area Population	Price
0	23086.80050	1.059034e+06
1	40173.07217	1.505891e+06
2	36882.15940	1.058988e+06
3	34310.24283	1.260617e+06
4	26354.10947	6.309435e+05


Q2. Find Features of Data-set and Number and Types .

```
print("Columns in the dataset:", df.columns.tolist())
print("Number of rows:", df.shape[0])
print("Number of columns:", df.shape[1])
print("Data types of each column:\n", df.dtypes)
```

Columns in the dataset: ['Area Income', 'Area House Age', 'Area No of Rooms', 'Area No of Bedrooms', 'Area Population', 'Price']
Number of rows: 5000
Number of columns: 6
Data types of each column:
Area Income float64
Area House Age float64
Area No of Rooms float64
Area No of Bedrooms float64
Area Population float64
Price float64
dtype: object


Q3.plot the USA housing histogram for price.

```
import matplotlib.pyplot as plt
plt.figure(figsize=(5, 5))
plt.hist(df['Price'], bins=30, color='skyblue', edgecolor='black')
plt.title('Distribution of House Prices')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```



Q4.plot Heat map of the correlation between each of columns.







```
import seaborn as sns
corr_matrix = df.corr()
plt.figure(figsize=(18, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Heatmap of Correlation Between Columns')
plt.show()
```




Q5. Plot Histograms and correlation scatter plots.

```
df.hist(bins=30, figsize=(15, 10), color='lightblue')
plt.suptitle('Histograms for All Features')
plt.show()
sns.pairplot(df)
plt.suptitle('Pairplot of All Features', y=1.02)
plt.show()
```

Histograms for All Features





Q6. Check Correlation of each feature with y

```
corr_with_price = df.corr()['Price'].sort_values(ascending=False)
print("Correlation of each feature with 'Price':\n", corr_with_price)
```

Correlation of each feature with 'Price':

Price	1.000000
Area Income	0.639734
Area House Age	0.452543
Area Population	0.408556
Area No of Rooms	0.335664

Area No of Bedrooms 0.171071
Name: Price, dtype: float64

Q7. Divide the dataset as 40% and 60% into test data and training data, respectively.

```
from sklearn.model_selection import train_test_split
X = df.drop('Price', axis=1)
y = df['Price']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.4, random_state=42)
print(f"Training set shape: {X_train.shape}, {y_train.shape}")
print(f"Testing set shape: {X_test.shape}, {y_test.shape}")
```

Training set shape: (3000, 5), (3000,)
Testing set shape: (2000, 5), (2000,)

Q8. Apply Linear Regression on Training Dataset and Calculate Coefficients

```
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)
print("Coefficients of Linear Regression model:\n", lr.coef_)
print("Intercept of Linear Regression model:", lr.intercept_)
```


Coefficients of Linear Regression model:
[2.15704132e+01 1.66552478e+05 1.19512534e+05 2.75895188e+03
1.52968610e+01]
Intercept of Linear Regression model: -2642239.251223582

Q9. Apply predict() Method to Calculate Predicted Value for Test Dataset

```
y_pred = lr.predict(X_test)
y_pred

array([1311074.59213156, 1239575.75693631, 1243916.2544329 , ...,
      840825.86579954, 1050491.02275121, 946720.46743172])

plt.figure(figsize=(5,5))
sns.scatterplot(x=y_test, y=y_pred, color='blue', alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
color='red', linestyle='--')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Scatter Plot: Actual Prices vs Predicted Prices')
plt.grid(True)
plt.show()
```

Q10. Calculate Error Metrics: MAE, MSE, and RMSE


```
from sklearn.metrics import mean_absolute_error, mean_squared_error
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)
print("Root Mean Squared Error (RMSE):", rmse)
```

```
Mean Absolute Error (MAE): 81331.22699046323
Mean Squared Error (MSE): 10119734875.098427
Root Mean Squared Error (RMSE): 100596.89296940749
```

Q11. Using distplot to Plot Predicted Values and Test Values for Correctness


```
plt.figure(figsize=(8, 5))
sns.scatterplot(x=y_test, y=y_pred, color='blue', alpha=0.6)
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Scatter Plot: Actual Prices vs Predicted Prices')
```

```
plt.grid(True)
plt.show()
```



Q12. Plot the Residual Histogram

```
residuals = y_test - y_pred
plt.figure(figsize=(8, 5))
sns.histplot(residuals, bins=30, kde=True, color='purple')
plt.title('Residual Histogram')
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.show()
```



Assignment 3

Q1. Select type and descriptions of dataset.

Q2. Take Logistic regression in the python library. Parameter tuning for optimizing classifier performance set for several parameter values namely: $C = 1.0$, $\lambda = 1.2 C$. The default value is 1.0. The smaller C value specifies the stronger regularization value. The Lambda parameter used to control the regularization strength. The increase in the lambda values the strength of the regularization effect.

Q3. Find confusion matrix using logistic regression without SGD.

Q4. Find ACCURACY, PRECISION, F-MEASURE, RECALL: LOGISTIC REGRESSION CLASSIFIER


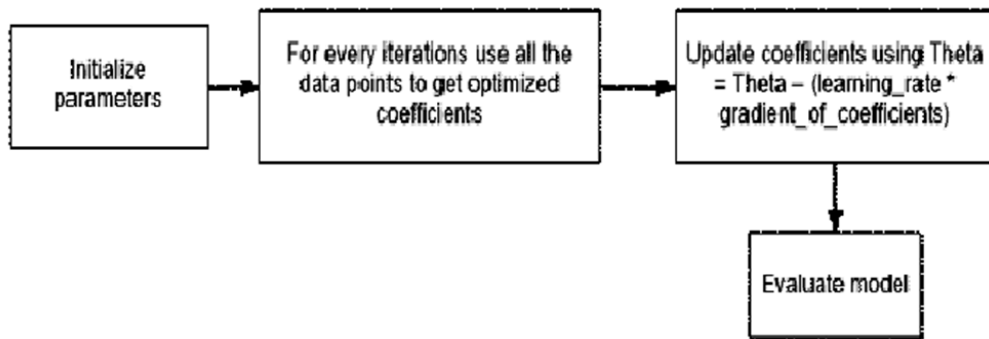
Q5. Learning rate that is derived from alpha and epsilon value. The learning rate managed the number of modification for estimating error. The alpha parameter is a constant value that multiplies the regularization term. The higher the value, the stronger the regularization. The default value is 0.0001. Epsilon parameter controlled the threshold value for getting precise prediction. Any distinction between the current prediction and the correct value is unobserved when this value less than the threshold value. The default value is 0.1. Experiment on this study set alpha value of 0.0001 and epsilon value of 0.1.

Q6. . Find confusion matrix using logistic regression with SGD.

Q7. . Find ACCURACY, PRECISION, F-MEASURE, RECALL: LOGISTIC REGRESSION CLASSIFIER with SGD.

Q8. Plot a graph Classifier result: Logistic regression and with Stochastic Gradient Descent for ACCURACY, PRECISION, F-MEASURE, RECALL.

Flowchart of the Assignment



```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score,
precision_score, recall_score, f1_score
from sklearn.linear_model import SGDClassifier
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
```

Q1. Load the cleaned kidney dataset(ckd)dataset and print its type and description

```
df = pd.read_excel('C:/Users/Asus/Desktop/DL Lab/DL Lab/Assignment
3/ckd_clean1.xlsx')
print(df.dtypes)
print(df.describe())
```

Age	int64
Blood Pressure	int64
Specific Gravity	float64
Albumin	int64
Sugar	int64
Red Blood Cells	object
Pus Cell	object
Pus Cell clumps	object
Bacteria	object
Blood Glucose Random	int64
Blood Urea	int64
Serum Creatinine	float64
Sodium	int64
Potassium	float64
Hemoglobin	float64
Packed Cell Volume	int64
White Blood Cell Count	int64
Red Blood Cell Count	float64
Hypertension	object
Diabetes Mellitus	object
Coronary Artery Disease	object
Appetite	object
Pedal Edema	object
Anemia	object
Class	int64

dtype: object

	Age	Blood Pressure	Specific Gravity	Albumin
Sugar \				
count	158.000000	158.000000	158.000000	158.000000
mean	49.563291	74.050633	1.019873	0.797468
std	15.512244	11.175381	0.005499	1.413130

```

0.813397
min      6.000000      50.000000      1.005000      0.000000
0.000000
25%      39.250000      60.000000      1.020000      0.000000
0.000000
50%      50.500000      80.000000      1.020000      0.000000
0.000000
75%      60.000000      80.000000      1.025000      1.000000
0.000000
max      83.000000      110.000000      1.025000      4.000000
5.000000

```

	Blood Glucose Random	Blood Urea	Serum Creatinine	Sodium
\ count	158.000000	158.000000	158.000000	158.000000
mean	131.341772	52.575949	2.188608	138.848101
std	64.939832	47.395382	3.077615	7.489421
min	70.000000	10.000000	0.400000	111.000000
25%	97.000000	26.000000	0.700000	135.000000
50%	115.500000	39.500000	1.100000	139.000000
75%	131.750000	49.750000	1.600000	144.000000
max	490.000000	309.000000	15.200000	150.000000

	Potassium	Hemoglobin	Packed Cell Volume	White Blood Cell
Count \				
count	158.000000	158.000000	158.000000	
158.000000				
mean	4.636709	13.687342	41.917722	
8475.949367				
std	3.476351	2.882204	9.105164	
3126.880181				
min	2.500000	3.100000	9.000000	
3800.000000				
25%	3.700000	12.600000	37.500000	
6525.000000				
50%	4.500000	14.250000	44.000000	
7800.000000				
75%	4.900000	15.775000	48.000000	
9775.000000				
max	47.000000	17.800000	54.000000	
26400.000000				

	Red Blood Cell Count	Class
count	158.000000	158.000000
mean	4.891772	0.272152
std	1.019364	0.446483
min	2.100000	0.000000
25%	4.500000	0.000000
50%	4.950000	0.000000
75%	5.600000	1.000000
max	8.000000	1.000000

Q2. Prepare the data for Logistic Regression Take Logistic regression in the python library. Parameter tuning for optimizing classifier performance set for several parameter values namely: $C=1.0$, $\lambda=1.2$ C . The default value is 1.0. The smaller C value specifies the stronger regularization value. The Lambda parameter used to control the regularization strength. The increase in the lambda values the strength of the regularization effect.

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
log_reg = LogisticRegression(max_iter=1000)
param_grid = {'C': [0.01, 0.1, 1.0, 10], 'penalty': ['l2']}
grid_search = GridSearchCV(log_reg, param_grid, scoring='accuracy',
cv=5)
grid_search.fit(X_train, y_train)
best_params = grid_search.best_params_
print("Best parameters: ", best_params)
categorical_cols = ['Red Blood Cells', 'Pus Cell', 'Pus Cell clumps',
'Bacteria', 'Hypertension', 'Diabetes Mellitus',
'Coronary Artery Disease', 'Appetite', 'Pedal
Edema', 'Anemia']
label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col].astype(str))
    label_encoders[col] = le
X = df.drop('Class', axis=1)
y = df['Class']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

Best parameters: {'C': 0.1, 'penalty': 'l2'}
```


Logistic Regression without SGD

```
lr = LogisticRegression(max_iter=1000, C=1.0)
lr.fit(X_train, y_train)

LogisticRegression(max_iter=1000)
```


Q3. Confusion Matrix

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
log_reg = LogisticRegression(C=best_params['C'], penalty='l2',
max_iter=1000)
log_reg.fit(X_train, y_train)
y_pred = log_reg.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



Q4. Accuracy, Precision, F-measure, Recall

```
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f_measure = f1_score(y_test, y_pred)
print("Without SGD - Accuracy: {}, Precision: {}, Recall: {}, F-
measure: {}".format(accuracy, precision, recall, f_measure))
```

Without SGD - Accuracy: 0.96875, Precision: 1.0, Recall: 0.8888888888888888, F-measure: 0.9411764705882353

Q5. Logistic Regression with SGD

```
from sklearn.linear_model import SGDClassifier
sgd_clf = SGDClassifier(loss='log', alpha=0.0001, epsilon=0.1,
max_iter=1000, tol=1e
sgd_clf.fit(X_train, y_train)

Cell In[20], line 2
    sgd_clf = SGDClassifier(loss='log', alpha=0.0001, epsilon=0.1,
max_iter=1000, tol=1e
^
SyntaxError: invalid syntax
```

Q6. Confusion Matrix with SGD

```
y_pred_sgd = sgd_clf.predict(X_test)
cm_sgd = confusion_matrix(y_test, y_pred_sgd)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

-----
-----
NameError                                Traceback (most recent call
last)
Cell In[19], line 1
----> 1 y_pred_sgd = sgd_clf.predict(X_test)
      2 cm_sgd = confusion_matrix(y_test, y_pred_sgd)
      3 plt.figure(figsize=(6, 4))

NameError: name 'sgd_clf' is not defined
```


Q7. Accuracy, Precision, F-measure, Recall with SGD

```
accuracy_sgd = accuracy_score(y_test, y_pred_sgd)
precision_sgd = precision_score(y_test, y_pred_sgd)
recall_sgd = recall_score(y_test, y_pred_sgd)
f_measure_sgd = f1_score(y_test, y_pred_sgd)
print("With SGD - Accuracy: {}, Precision: {}, Recall: {}, F-measure:
{}".format(accuracy_sgd, precision_sgd, recall_sgd, f_measure_sgd))
```

With SGD - Accuracy: 0.28125, Precision: 0.28125, Recall: 1.0, F-measure: 0.43902439024390244

Q8. Plotting the results

```
labels = ['Accuracy', 'Precision', 'Recall', 'F-measure']
without_sgd = [accuracy, precision, recall, f_measure]
with_sgd = [accuracy_sgd, precision_sgd, recall_sgd, f_measure_sgd]
x = range(len(labels))
plt.plot(x, without_sgd, label='Without SGD', marker='o')
plt.plot(x, with_sgd, label='With SGD', marker='o')
plt.xticks(x, labels)
plt.ylabel('Score')
plt.title('Comparison of Logistic Regression and SGD')
plt.legend()
plt.show()
```



Vishal Marandi |

BTECH/10119/21 |

Deep Learning Assignment - 3

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score,
precision_score, recall_score, f1_score
from sklearn.linear_model import SGDClassifier
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
```

Q1. Load the dataset and print its type and description

```
df = pd.read_csv('/content/Thyroid.txt')
print(df.dtypes)
print(df.describe())
```

on_thyroxine	int64
query_on_thyroxine	int64
on_antithyroid_medication	int64
thyroid_surgery	int64
query_hypothyroid	int64
query_hyperthyroid	int64
pregnant	int64
sick	int64
tumor	int64
lithium	int64
goitre	int64
TSH_measured	int64
T3_measured	int64
TT4_measured	int64
T4U_measured	int64
FTI_measured	int64
age	int64
sex	int64
TSH	float64
T3	float64
TT4	float64
T4U	float64
FTI	float64
classes	int64

```

dtype: object
count    on_thyroxine    query_on_thyroxine    on_antithyroid_medication \
mean      0.145305      0.017449      0.013325
std       0.352464      0.130959      0.114680
min       0.000000      0.000000      0.000000
25%       0.000000      0.000000      0.000000
50%       0.000000      0.000000      0.000000
75%       0.000000      0.000000      0.000000
max       1.000000      1.000000      1.000000

```

```

thyroid_surgery    query_hypothyroid    query_hyperthyroid
pregnant \
count    3152.000000    3152.000000    3152.000000
3152.000000
mean      0.032995      0.075825      0.077094
0.019987
std       0.178652      0.264760      0.266783
0.139979
min       0.000000      0.000000      0.000000
0.000000
25%       0.000000      0.000000      0.000000
0.000000
50%       0.000000      0.000000      0.000000
0.000000
75%       0.000000      0.000000      0.000000
0.000000
max       1.000000      1.000000      1.000000
1.000000

```

```

sick    tumor    lithium    ...    T4U_measured
FTI_measured \
count    3152.000000    3152.000000    3152.000000    ...    3152.000000
3152.000000
mean      0.030774      0.012690      0.000635    ...      0.921320
0.921637
std       0.172733      0.111952      0.025186    ...      0.269282
0.268785
min       0.000000      0.000000      0.000000    ...      0.000000
0.000000
25%       0.000000      0.000000      0.000000    ...      1.000000
1.000000
50%       0.000000      0.000000      0.000000    ...      1.000000
1.000000
75%       0.000000      0.000000      0.000000    ...      1.000000
1.000000
max       1.000000      1.000000      1.000000    ...      1.000000
1.000000

```

```

age    sex    TSH    T3    TT4

```

\	count	3152.000000	3152.000000	3152.000000	3152.000000	3152.000000
	mean	51.359772	0.295051	5.689591	1.992291	109.016783
	std	19.226176	0.456138	22.953469	1.036009	45.085350
	min	1.000000	0.000000	0.000000	0.000000	2.000000
	25%	35.000000	0.000000	0.000000	1.400000	83.000000
	50%	54.000000	0.000000	0.700000	1.900000	104.000000
	75%	67.000000	1.000000	2.200000	2.400000	128.000000
	max	98.000000	1.000000	530.000000	10.200000	450.000000

	T4U	FTI	classes
count	3152.000000	3152.000000	3152.000000
mean	0.983284	114.809613	0.091371
std	0.234424	59.028114	0.288181
min	0.000000	0.000000	0.000000
25%	0.850000	90.000000	0.000000
50%	0.960000	107.000000	0.000000
75%	1.070000	128.000000	0.000000
max	2.210000	881.000000	1.000000

[8 rows x 24 columns]

Q2. Prepare the data for Logistic Regression

```
X = df.drop('classes', axis=1)
y = df['classes']
# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Logistic Regression without SGD

```
lr = LogisticRegression(max_iter=1000, C=1.0)
lr.fit(X_train, y_train)
```

```
LogisticRegression(max_iter=1000)
```

Q3. Confusion Matrix

```
y_pred = lr.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix (without SGD):\n", cm)
```

```
Confusion Matrix (without SGD):
[[552  10]
 [ 23  46]]
```

Q4. Accuracy, Precision, F-measure, Recall

```
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f_measure = f1_score(y_test, y_pred)
print("Without SGD - Accuracy: {}, Precision: {}, Recall: {}, F-
measure: {}".format(accuracy, precision, recall, f_measure))
```

```
Without SGD - Accuracy: 0.9477020602218701, Precision:
0.8214285714285714, Recall: 0.6666666666666666, F-measure:
0.7359999999999999
```

Q5. Logistic Regression with SGD

```
sgd = SGDClassifier(loss='log_loss', alpha=0.0001, max_iter=1000,
tol=1e-3, random_state=42)
sgd.fit(X_train, y_train)
```

```
SGDClassifier(loss='log_loss', random_state=42)
```

Q6. Confusion Matrix with SGD

```
y_pred_sgd = sgd.predict(X_test)
cm_sgd = confusion_matrix(y_test, y_pred_sgd)
print("Confusion Matrix (with SGD):\n", cm_sgd)
```

```
Confusion Matrix (with SGD):
[[562   0]
 [ 69   0]]
```

Q7. Accuracy, Precision, F-measure, Recall with SGD

```
accuracy_sgd = accuracy_score(y_test, y_pred_sgd)
precision_sgd = precision_score(y_test, y_pred_sgd)
recall_sgd = recall_score(y_test, y_pred_sgd)
f_measure_sgd = f1_score(y_test, y_pred_sgd)
print("With SGD - Accuracy: {}, Precision: {}, Recall: {}, F-measure: {}".format(accuracy_sgd, precision_sgd, recall_sgd, f_measure_sgd))
```


```
With SGD - Accuracy: 0.8906497622820919, Precision: 0.0, Recall: 0.0, F-measure: 0.0
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

Q8. Plotting the results

```
labels = ['Accuracy', 'Precision', 'Recall', 'F-measure']
without_sgd = [accuracy, precision, recall, f_measure]
with_sgd = [accuracy_sgd, precision_sgd, recall_sgd, f_measure_sgd]
x = range(len(labels))
plt.plot(x, without_sgd, label='Without SGD', marker='o')
plt.plot(x, with_sgd, label='With SGD', marker='o')
plt.xticks(x, labels)
plt.ylabel('Score')
plt.title('Comparison of Logistic Regression and SGD')
plt.legend()
plt.show()
```


Comparison of Logistic Regression and SGD



[illegible]

```

9      float64
10     float64
11     float64
12     float64
13     float64
14      int64
dtype: object

```

	1	2	3	4	5
6 \					
count	270.000000	270.000000	270.000000	270.000000	270.000000
mean	54.433333	0.677778	3.174074	131.344444	249.659259
std	9.109067	0.468195	0.950090	17.861608	51.686237
min	29.000000	0.000000	1.000000	94.000000	126.000000
25%	48.000000	0.000000	3.000000	120.000000	213.000000
50%	55.000000	1.000000	3.000000	130.000000	245.000000
75%	61.000000	1.000000	4.000000	140.000000	280.000000
max	77.000000	1.000000	4.000000	200.000000	564.000000

	7	8	9	10	11
12 \					
count	270.000000	270.000000	270.000000	270.000000	270.000000
mean	1.022222	149.677778	0.329630	1.050000	1.585185
std	0.997891	23.165717	0.470952	1.145210	0.614390
min	0.000000	71.000000	0.000000	0.000000	1.000000
25%	0.000000	133.000000	0.000000	0.000000	1.000000
50%	2.000000	153.500000	0.000000	0.800000	2.000000
75%	2.000000	166.000000	1.000000	1.600000	2.000000
max	2.000000	202.000000	1.000000	6.200000	3.000000

	13	14
count	270.000000	270.000000
mean	4.696296	1.444444
std	1.940659	0.497827

min	3.000000	1.000000
25%	3.000000	1.000000
50%	3.000000	1.000000
75%	7.000000	2.000000
max	7.000000	2.000000

Q2. Prepare the data for Logistic Regression

```
X = df.iloc[:, :-1] # All columns except the last one
y = df.iloc[:, -1] # The last column as target
# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Logistic Regression without SGD

```
lr = LogisticRegression(max_iter=1000, C=1.0)
lr.fit(X_train, y_train)

LogisticRegression(max_iter=1000)
```

Q3. Confusion Matrix

```
y_pred = lr.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix (without SGD):\n", cm)

Confusion Matrix (without SGD):
[[32  1]
 [ 3 18]]
```

Q4. Accuracy, Precision, F-measure, Recall

```
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f_measure = f1_score(y_test, y_pred)
print("Without SGD - Accuracy: {}, Precision: {}, Recall: {}, F-
measure: {}".format(accuracy, precision, recall, f_measure))

Without SGD - Accuracy: 0.9259259259259259, Precision:
0.9142857142857143, Recall: 0.9696969696969697, F-measure:
0.9411764705882354
```

Q5. Logistic Regression with SGD

```
sgd = SGDClassifier(loss='log_loss', alpha=0.0001, max_iter=1000,  
tol=1e-3, random_state=42)  
sgd.fit(X_train, y_train)
```

```
SGDClassifier(loss='log_loss', random_state=42)
```

Q6. Confusion Matrix with SGD

```
y_pred_sgd = sgd.predict(X_test)  
cm_sgd = confusion_matrix(y_test, y_pred_sgd)  
print("Confusion Matrix (with SGD):\n", cm_sgd)
```

```
Confusion Matrix (with SGD):  
[[28  5]  
 [ 5 16]]
```

Q7. Accuracy, Precision, F-measure, Recall with SGD


```
accuracy_sgd = accuracy_score(y_test, y_pred_sgd)  
precision_sgd = precision_score(y_test, y_pred_sgd)  
recall_sgd = recall_score(y_test, y_pred_sgd)  
f_measure_sgd = f1_score(y_test, y_pred_sgd)  
print("With SGD - Accuracy: {}, Precision: {}, Recall: {}, F-measure:  
{ }".format(accuracy_sgd, precision_sgd, recall_sgd, f_measure_sgd))
```

```
With SGD - Accuracy: 0.8148148148148148, Precision:  
0.8484848484848485, Recall: 0.8484848484848485, F-measure:  
0.8484848484848486
```

Q8. Plotting the results

```
labels = ['Accuracy', 'Precision', 'Recall', 'F-measure']  
without_sgd = [accuracy, precision, recall, f_measure]  
with_sgd = [accuracy_sgd, precision_sgd, recall_sgd, f_measure_sgd]  
x = range(len(labels))  
plt.plot(x, without_sgd, label='Without SGD', marker='o')  
plt.plot(x, with_sgd, label='With SGD', marker='o')  
plt.xticks(x, labels)  
plt.ylabel('Score')  
plt.title('Comparison of Logistic Regression and SGD')
```

```
plt.legend()  
plt.show()
```



Assignment 4

- Q1. Load load_breast_cancer() built in dataset.**
- Q2. convert dataset to pandas dataframe.**
- Q3. Append dataframe containing tumor features with diagnostic outcomes.**
These labels will be used for supervised learning.
- Q4. Find relation in diagnosis by using mean.**
- Q5. create dataframes - one for positive, one for negative(diagnosis).**
- Q6. Visualize tumor characteristics for positive and negatives diagnoses**
- Q7. Create 'for loop' to iterate through tumor features and compare with
histograms(20:-1)**
- Q8. Find Visualization of relationships between features and diagnoses.**
- Q9. Split data into testing and training sets. Use 80% for training**
- Q10. normalize features to account for feature scaling**
- Q11. Apply DT, NB, KNN Logistic Regression and SVM to find accuracy,
F1-Score and Recall.**
- Q12. Further apply the PCA to do the same exercise in Q11.**
- Q13. Do the Comparative study with PCA and WITHOUT PCA.**

Q1. Load `load_breast_cancer()` built in dataset.

```
import pandas as pd
import numpy as np
from sklearn.datasets import load_breast_cancer
breast_cancer_data = load_breast_cancer()
print(breast_cancer_data)

{'data': array([[1.799e+01, 1.038e+01, 1.228e+02, ..., 2.654e-01,
4.601e-01,
1.189e-01],
[2.057e+01, 1.777e+01, 1.329e+02, ..., 1.860e-01, 2.750e-01,
8.902e-02],
[1.969e+01, 2.125e+01, 1.300e+02, ..., 2.430e-01, 3.613e-01,
8.758e-02],
...,
[1.660e+01, 2.808e+01, 1.083e+02, ..., 1.418e-01, 2.218e-01,
7.820e-02],
[2.060e+01, 2.933e+01, 1.401e+02, ..., 2.650e-01, 4.087e-01,
1.240e-01],
[7.760e+00, 2.454e+01, 4.792e+01, ..., 0.000e+00, 2.871e-01,
7.039e-02]]), 'target': array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0,
0,
1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0,
0,
1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0,
1,
1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1,
0,
0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
1,
1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1,
1,
1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0,
0,
0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0,
0,
1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1,
1,
1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1,
1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1,
```



```

0,      1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0,
0,      0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
0,      0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0,
0,      1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1,
1,      1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1,
0,      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1,
1,      1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0,
0,      1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1,
1,      1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1,
1,      1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1,
1,      1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1]),
'frame': None, 'target_names': array(['malignant', 'benign'],
dtype='<U9'), 'DESCR': '.. _breast_cancer_dataset:\n\nBreast cancer
wisconsin (diagnostic) dataset\
n-----\n\n**Data Set
Characteristics:**\n\n:Number of Instances: 569\n\n:Number of
Attributes: 30 numeric, predictive attributes and the class\n\n
n:Attribute Information:\n      - radius (mean of distances from center
to points on the perimeter)\n      - texture (standard deviation of
gray-scale values)\n      - perimeter\n      - area\n      - smoothness
(local variation in radius lengths)\n      - compactness (perimeter^2 /
area - 1.0)\n      - concavity (severity of concave portions of the
contour)\n      - concave points (number of concave portions of the
contour)\n      - symmetry\n      - fractal dimension ("coastline
approximation" - 1)\n\n      The mean, standard error, and "worst" or
largest (mean of the three\n      worst/largest values) of these
features were computed for each image,\n      resulting in 30 features.
For instance, field 0 is Mean Radius, field\n      10 is Radius SE,
field 20 is Worst Radius.\n\n      - class:\n      - WDBC-
Malignant\n      - WDBC-Benign\n\n:Summary Statistics:\n\n
n===== \n
Min      Max\n===== \n
nradius (mean):                6.981  28.11\ntexture (mean):
9.71  39.28\nnperimeter (mean):                43.79  188.5\nnarea
(mean):                143.5  2501.0\nnsmoothness (mean):
0.053  0.163\nncompactness (mean):                0.019  0.345\
nconcavity (mean):                0.0  0.427\nnconcave points

```

```

(mean):          0.0    0.201\nsymmetry (mean):
0.106  0.304\nfractal dimension (mean):          0.05    0.097\
nradius (standard error):          0.112  2.873\ntexture (standard
error):          0.36    4.885\nperimeter (standard error):
0.757  21.98\narea (standard error):          6.802  542.2\
nsmoothness (standard error):          0.002  0.031\ncompactness
(standard error):          0.002  0.135\nconcavity (standard error):
0.0    0.396\nconcave points (standard error):          0.0    0.053\
nsymmetry (standard error):          0.008  0.079\nfractal dimension
(standard error):  0.001  0.03\nradius (worst):
7.93   36.04\ntexture (worst):          12.02  49.54\
nperimeter (worst):          50.41  251.2\narea (worst):
185.2  4254.0\nsmoothness (worst):          0.071  0.223\
ncompactness (worst):          0.027  1.058\nconcavity
(worst):          0.0    1.252\nconcave points (worst):
0.0    0.291\nsymmetry (worst):          0.156  0.664\
nfractal dimension (worst):          0.055  0.208\
n===== \n\nMissing
Attribute Values: None\n\nClass Distribution: 212 - Malignant, 357 -
Benign\n\nCreator: Dr. William H. Wolberg, W. Nick Street, Olvi L.
Mangasarian\n\nDonor: Nick Street\n\nDate: November, 1995\n\nThis is
a copy of UCI ML Breast Cancer Wisconsin (Diagnostic) datasets.\n
nhttps://goo.gl/U2Uwz2\n\nFeatures are computed from a digitized image
of a fine needle\naspirate (FNA) of a breast mass. They describe\
ncharacteristics of the cell nuclei present in the image.\n\
nSeparating plane described above was obtained using\nMultisurface
Method-Tree (MSM-T) [K. P. Bennett, "Decision Tree\nConstruction Via
Linear Programming." Proceedings of the 4th\nMidwest Artificial
Intelligence and Cognitive Science Society,\npp. 97-101, 1992], a
classification method which uses linear\nprogramming to construct a
decision tree. Relevant features\nwere selected using an exhaustive
search in the space of 1-4\nfeatures and 1-3 separating planes.\n\nThe
actual linear program used to obtain the separating plane\nin the 3-
dimensional space is that described in:\n[K. P. Bennett and O. L.
Mangasarian: "Robust Linear\nProgramming Discrimination of Two
Linearly Inseparable Sets",\nOptimization Methods and Software 1,
1992, 23-34].\n\nThis database is also available through the UW CS ftp
server:\n\nftp ftp.cs.wisc.edu\ncd math-prog/cpo-dataset/machine-
learn/WDBC/\n\n.. dropdown:: References\n\n - W.N. Street, W.H.
Wolberg and O.L. Mangasarian. Nuclear feature extraction\n for
breast tumor diagnosis. IS&T/SPIE 1993 International Symposium on\n
Electronic Imaging: Science and Technology, volume 1905, pages 861-
870,\n San Jose, CA, 1993.\n - O.L. Mangasarian, W.N. Street and
W.H. Wolberg. Breast cancer diagnosis and\n prognosis via linear
programming. Operations Research, 43(4), pages 570-577,\n July-
August 1995.\n - W.H. Wolberg, W.N. Street, and O.L. Mangasarian.
Machine learning techniques\n to diagnose breast cancer from fine-
needle aspirates. Cancer Letters 77 (1994)\n 163-171.\n',
'feature_names': array(['mean radius', 'mean texture', 'mean

```

```

perimeter', 'mean area',
    'mean smoothness', 'mean compactness', 'mean concavity',
    'mean concave points', 'mean symmetry', 'mean fractal
dimension',
    'radius error', 'texture error', 'perimeter error', 'area
error',
    'smoothness error', 'compactness error', 'concavity error',
    'concave points error', 'symmetry error',
    'fractal dimension error', 'worst radius', 'worst texture',
    'worst perimeter', 'worst area', 'worst smoothness',
    'worst compactness', 'worst concavity', 'worst concave points',
    'worst symmetry', 'worst fractal dimension'], dtype='<U23'),
'filename': 'breast_cancer.csv', 'data_module':
'sklearn.datasets.data'}

X = breast_cancer_data.data
y = breast_cancer_data.target
feature_names = breast_cancer_data.feature_names
print(feature_names)

['mean radius' 'mean texture' 'mean perimeter' 'mean area'
'mean smoothness' 'mean compactness' 'mean concavity'
'mean concave points' 'mean symmetry' 'mean fractal dimension'
'radius error' 'texture error' 'perimeter error' 'area error'
'smoothness error' 'compactness error' 'concavity error'
'concave points error' 'symmetry error' 'fractal dimension error'
'worst radius' 'worst texture' 'worst perimeter' 'worst area'
'worst smoothness' 'worst compactness' 'worst concavity'
'worst concave points' 'worst symmetry' 'worst fractal dimension']

```

Q2. convert dataset to pandas dataframe.

```

df_features = pd.DataFrame(X, columns=feature_names)
print(df_features.head())

```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness \
0	17.99	10.38	122.80	1001.0	0.11840
1	20.57	17.77	132.90	1326.0	0.08474
2	19.69	21.25	130.00	1203.0	0.10960
3	11.42	20.38	77.58	386.1	0.14250
4	20.29	14.34	135.10	1297.0	0.10030

	mean compactness	mean concavity	mean concave points	mean symmetry \
--	------------------	----------------	---------------------	-----------------

0	0.27760	0.3001	0.14710
0.2419			
1	0.07864	0.0869	0.07017
0.1812			
2	0.15990	0.1974	0.12790
0.2069			
3	0.28390	0.2414	0.10520
0.2597			
4	0.13280	0.1980	0.10430
0.1809			

	mean fractal dimension	...	worst radius	worst texture	worst perimeter \
0	0.07871	...	25.38	17.33	
184.60					
1	0.05667	...	24.99	23.41	
158.80					
2	0.05999	...	23.57	25.53	
152.50					
3	0.09744	...	14.91	26.50	
98.87					
4	0.05883	...	22.54	16.67	
152.20					

	worst area	worst smoothness	worst compactness	worst concavity \
0	2019.0	0.1622	0.6656	0.7119
1	1956.0	0.1238	0.1866	0.2416
2	1709.0	0.1444	0.4245	0.4504
3	567.7	0.2098	0.8663	0.6869
4	1575.0	0.1374	0.2050	0.4000

	worst concave points	worst symmetry	worst fractal dimension
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678

[5 rows x 30 columns]

Q3. Append dataframe containing tumor features with diagnostic outcomes.

```
df_target = pd.Series(y, name='diagnosis')
df = pd.concat([df_features, df_target], axis=1)
print(df.head())
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness \
0	17.99	10.38	122.80	1001.0	
0.11840					

1	20.57	17.77	132.90	1326.0
0.08474				
2	19.69	21.25	130.00	1203.0
0.10960				
3	11.42	20.38	77.58	386.1
0.14250				
4	20.29	14.34	135.10	1297.0
0.10030				
mean compactness mean concavity mean concave points mean				
symmetry \				
0	0.27760	0.3001	0.14710	
0.2419				
1	0.07864	0.0869	0.07017	
0.1812				
2	0.15990	0.1974	0.12790	
0.2069				
3	0.28390	0.2414	0.10520	
0.2597				
4	0.13280	0.1980	0.10430	
0.1809				
mean fractal dimension ... worst texture worst perimeter worst				
area \				
0	0.07871	...	17.33	184.60
2019.0				
1	0.05667	...	23.41	158.80
1956.0				
2	0.05999	...	25.53	152.50
1709.0				
3	0.09744	...	26.50	98.87
567.7				
4	0.05883	...	16.67	152.20
1575.0				
worst smoothness worst compactness worst concavity worst concave				
points \				
0	0.1622	0.6656	0.7119	
0.2654				
1	0.1238	0.1866	0.2416	
0.1860				
2	0.1444	0.4245	0.4504	
0.2430				
3	0.2098	0.8663	0.6869	
0.2575				
4	0.1374	0.2050	0.4000	
0.1625				
worst symmetry worst fractal dimension diagnosis				
0	0.4601	0.11890	0	

1	0.2750	0.08902	0
2	0.3613	0.08758	0
3	0.6638	0.17300	0
4	0.2364	0.07678	0

[5 rows x 31 columns]

Q4.Find relation in diagnosis by using mean.

```
mean_features = df.groupby('diagnosis').mean()
print(mean_features)
```

	mean radius	mean texture	mean perimeter	mean area \
diagnosis				
0	17.462830	21.604906	115.365377	978.376415
1	12.146524	17.914762	78.075406	462.790196

	mean smoothness	mean compactness	mean concavity \
diagnosis			
0	0.102898	0.145188	0.160775
1	0.092478	0.080085	0.046058

	mean concave points	mean symmetry	mean fractal dimension
...			
diagnosis			
...			
0	0.087990	0.192909	0.062680
...			
1	0.025717	0.174186	0.062867
...			

	worst radius	worst texture	worst perimeter	worst
area \				
diagnosis				
0	21.134811	29.318208	141.370330	1422.286321
1	13.379801	23.515070	87.005938	558.899440

	worst smoothness	worst compactness	worst concavity \
diagnosis			
0	0.144845	0.374824	0.450606
1	0.124959	0.182673	0.166238

	worst concave points	worst symmetry	worst fractal
dimension			
diagnosis			
0	0.182237	0.323468	

```
0.091530
1          0.074444          0.270246
0.079442

[2 rows x 30 columns]
```

Q5. create to dataframes - one for positive, one for negative(diagnosis).

```
df_positive = df[df['diagnosis'] == 1]
df_negative = df[df['diagnosis'] == 0]
print("Positive diagnosis (Malignant):")
print(df_positive.head())
print("\nNegative diagnosis (Benign):")
print(df_negative.head())
```

Positive diagnosis (Malignant):

	mean radius	mean texture	mean perimeter	mean area	mean smoothness \
19	13.540	14.36	87.46	566.3	0.09779
20	13.080	15.71	85.63	520.0	0.10750
21	9.504	12.44	60.34	273.9	0.10240
37	13.030	18.42	82.61	523.8	0.08983
46	8.196	16.84	51.71	201.9	0.08600

	mean compactness	mean concavity	mean concave points	mean symmetry \
19	0.08129	0.06664	0.047810	0.1885
20	0.12700	0.04568	0.031100	0.1967
21	0.06492	0.02956	0.020760	0.1815
37	0.03766	0.02562	0.029230	0.1467
46	0.05943	0.01588	0.005917	0.1769

	mean fractal dimension	...	worst texture	worst perimeter	worst area \
19	0.05766	...	19.26	99.70	711.2
20	0.06811	...	20.49	96.09	630.5
21	0.06905	...	15.66	65.13	

314.9				
37	0.05863	...	22.81	84.46
545.9				
46	0.06503	...	21.96	57.26
242.2				

	worst smoothness	worst compactness	worst concavity	\
19	0.14400	0.17730	0.23900	
20	0.13120	0.27760	0.18900	
21	0.13240	0.11480	0.08867	
37	0.09701	0.04619	0.04833	
46	0.12970	0.13570	0.06880	

	worst concave points	worst symmetry	worst fractal dimension
diagnosis			
19	0.12880	0.2977	0.07259
1			
20	0.07283	0.3184	0.08183
1			
21	0.06227	0.2450	0.07773
1			
37	0.05013	0.1987	0.06169
1			
46	0.02564	0.3105	0.07409
1			

[5 rows x 31 columns]

Negative diagnosis (Benign):

	mean radius	mean texture	mean perimeter	mean area	mean smoothness
\					
0	17.99	10.38	122.80	1001.0	0.11840
1	20.57	17.77	132.90	1326.0	0.08474
2	19.69	21.25	130.00	1203.0	0.10960
3	11.42	20.38	77.58	386.1	0.14250
4	20.29	14.34	135.10	1297.0	0.10030

	mean compactness	mean concavity	mean concave points	mean symmetry
\				
0	0.27760	0.3001	0.14710	0.2419
1	0.07864	0.0869	0.07017	0.1812
2	0.15990	0.1974	0.12790	0.2069


```

3          0.28390          0.2414          0.10520
0.2597
4          0.13280          0.1980          0.10430
0.1809

    mean fractal dimension ... worst texture worst perimeter worst
area \
0          0.07871 ...          17.33          184.60
2019.0
1          0.05667 ...          23.41          158.80
1956.0
2          0.05999 ...          25.53          152.50
1709.0
3          0.09744 ...          26.50          98.87
567.7
4          0.05883 ...          16.67          152.20
1575.0

    worst smoothness worst compactness worst concavity worst concave
points \
0          0.1622          0.6656          0.7119
0.2654
1          0.1238          0.1866          0.2416
0.1860
2          0.1444          0.4245          0.4504
0.2430
3          0.2098          0.8663          0.6869
0.2575
4          0.1374          0.2050          0.4000
0.1625

    worst symmetry worst fractal dimension diagnosis
0          0.4601          0.11890          0
1          0.2750          0.08902          0
2          0.3613          0.08758          0
3          0.6638          0.17300          0
4          0.2364          0.07678          0

[5 rows x 31 columns]

```


Q6. Visualize tumor characteristics for positive and negatives diagnoses

```


import matplotlib.pyplot as plt
df_positive.hist(figsize=(20, 20), bins=20, color='red', alpha=0.5)
plt.suptitle('Positive Diagnosis (Malignant)')
plt.show()
df_negative.hist(figsize=(20, 20), bins=20, color='blue', alpha=0.5)
plt.suptitle('Negative Diagnosis (Benign)')
plt.show()

```

Positive Diagnosis (Malignant)



Negative Diagnosis (Benign)





Q7. Create 'for loop' to enerate though tumor features and compare with histograms(20:-1)



```
features = df_features.columns
for feature in features:
    plt.figure(figsize=(10, 5))
    plt.hist(df_positive[feature], bins=20, alpha=0.5, label='Positive (Malignant)', color='red')
    plt.hist(df_negative[feature], bins=20, alpha=0.5, label='Negative
```



```



(Benign)', color='blue')
plt.title(f'Histogram of {feature}')
plt.xlabel(feature)
plt.ylabel('Frequency')
plt.legend()
plt.show()



```
































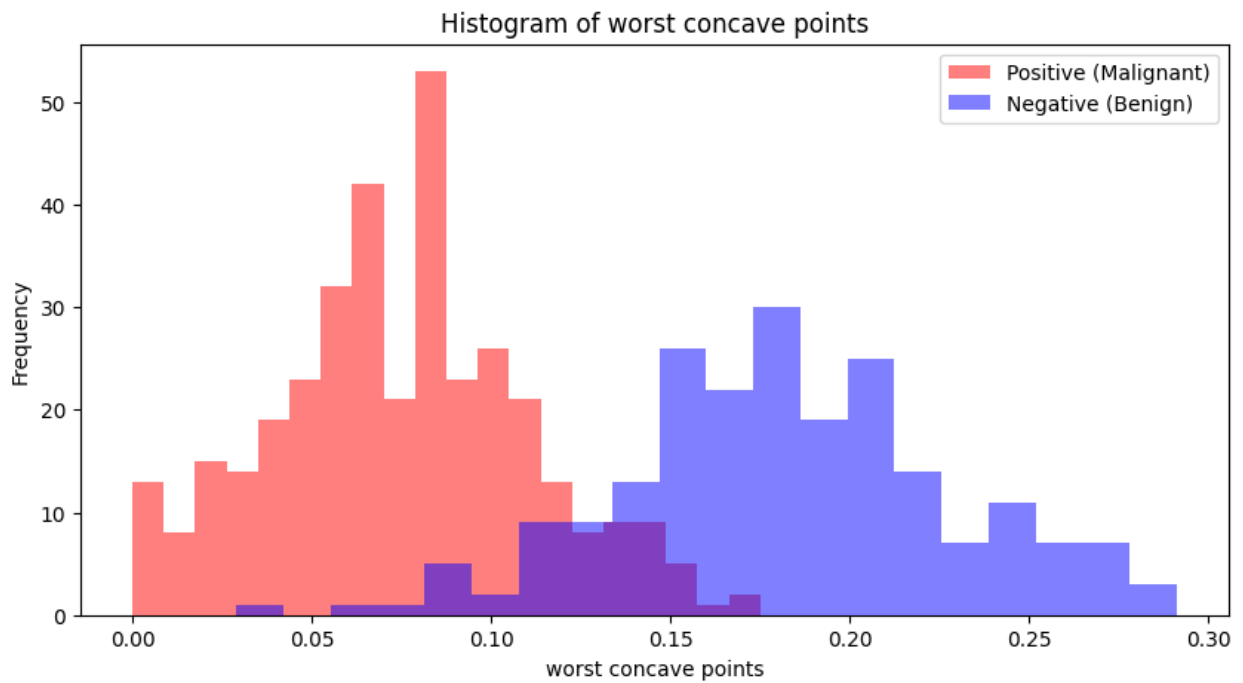


















Q8. Find Visualization of relationships between features and diagnoses.

```
import seaborn as sns
sns.pairplot(df, hue='diagnosis', vars=features[:5])
plt.show()
```



Q9. Split data into testing and training set. Use 80% for training

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)
```

Training set shape: (455, 30)
Testing set shape: (114, 30)

Q10. normalize features to account for feature scaling

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Q11. Apply DT, NB, KNN Logistic Regression and SVM to find accuracy, F1-Score and Recall.

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, f1_score, recall_score

models = {
    'Decision Tree': DecisionTreeClassifier(),
    'Naive Bayes': GaussianNB(),
    'K-Nearest Neighbors': KNeighborsClassifier(),
    'Logistic Regression': LogisticRegression(max_iter=10000),
    'Support Vector Machine': SVC()
}

results = {}
for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)
    results[name] = {
        'Accuracy': accuracy_score(y_test, y_pred),
        'F1 Score': f1_score(y_test, y_pred),
        'Recall': recall_score(y_test, y_pred)
    }

for name, metrics in results.items():
    print(f"{name}:")
    for metric, value in metrics.items():
        print(f"    {metric}: {value:.4f}")
```

```
C:\Users\Asus\AppData\Local\Programs\Python\Python39\lib\site-
packages\joblib\externals\loky\backend\context.py:136: UserWarning:
Could not find the number of physical cores for the following reason:
[WinError 2] The system cannot find the file specified
Returning the number of logical cores instead. You can silence this
warning by setting LOKY_MAX_CPU_COUNT to the number of cores you want
to use.
```

```
warnings.warn(
    File "C:\Users\Asus\AppData\Local\Programs\Python\Python39\lib\site-
packages\joblib\externals\loky\backend\context.py", line 257, in
_count_physical_cores
    cpu_info = subprocess.run(
    File "C:\Users\Asus\AppData\Local\Programs\Python\Python39\lib\
subprocess.py", line 505, in run
    with Popen(*popenargs, **kwargs) as process:
    File "C:\Users\Asus\AppData\Local\Programs\Python\Python39\lib\
subprocess.py", line 951, in __init__
```

```
self._execute_child(args, executable, preexec_fn, close_fds,  
File "C:\Users\Asus\AppData\Local\Programs\Python\Python39\lib\  
subprocess.py", line 1420, in _execute_child  
hp, ht, pid, tid = _winapi.CreateProcess(executable, args,
```

Decision Tree:

Accuracy: 0.9386

F1 Score: 0.9504

Recall: 0.9437

Naive Bayes:

Accuracy: 0.9649

F1 Score: 0.9722

Recall: 0.9859

K-Nearest Neighbors:

Accuracy: 0.9474

F1 Score: 0.9577

Recall: 0.9577

Logistic Regression:

Accuracy: 0.9737

F1 Score: 0.9790

Recall: 0.9859

Support Vector Machine:

Accuracy: 0.9825

F1 Score: 0.9861

Recall: 1.0000

Q12. Further apply the PCA to do the same exercise in Q11.

Q13. Do the Comprative study **with** PCA **and** WITHout PCA.

Part 4.2

Q1. Apply Mc-Culloch Pitts model for the following Boolean operations:

Plot the decision boundaries for 2 - dimensional and 3-dimensional.

Implement AND with 3-variables and threshold 3.

Implement OR with 3-variables and threshold 1.

Implement X_1 and $\neg X_2$ with 2-variables and threshold 1.

Implement NOR with 2-variables and threshold 0.

Implement NOT with 2-variables and threshold 0.

Implement OR with 2-variables and threshold 1.


Implement AND with 2-variables and threshold 2.

Implement OR with 3-variables and threshold 1.

Q1. Apply Mc-Culloch Pits model for the following Boolean operations: Plot the decision boundaries for 2 - dimensional and 3-dimensional. Implement AND with 3-variables and threshold 3.

```
import numpy as np
import matplotlib.pyplot as plt
fig = plt.figure(figsize=(10, 7))
ax = fig.add_subplot(111, projection='3d')
x1 = np.linspace(0, 3, 100)
x2 = np.linspace(0, 3, 100)
X1, X2 = np.meshgrid(x1, x2)
X3 = 3 - X1 - X2
ax.plot_surface(X1, X2, X3, alpha=0.5)
points = np.array([(i, j, k) for i in range(4) for j in range(4) for k
in range(4) if i + j + k >= 3])
ax.scatter(points[:, 0], points[:, 1], points[:, 2], color='r', s=50)
for point in points:
    ax.text(*point, f'({point[0]},{point[1]},{point[2]})',
color='black')
ax.set_xlabel('X1')
ax.set_ylabel('X2')
ax.set_zlabel('X3')
ax.set_title('3D AND Operation with Threshold 3')
plt.show()
```

3D AND Operation with Threshold 3




Implement OR with 3-variables and threshold 1.

```
x = np.linspace(0, 1, 100)
X1, X2 = np.meshgrid(x, x)
fig = plt.figure(figsize=(12, 8))
ax = fig.add_subplot(111, projection='3d')
x3_values = [0, 0.5, 1]
for x3 in x3_values:
    Z = (X1 + X2) >= (1 - x3)
    ax.plot_surface(X1, X2, x3 * np.ones_like(X1),
                    facecolors=plt.cm.gray(Z), alpha=0.5,
                    edgecolor='none')
ax.set_xlabel('$x_1$')
ax.set_ylabel('$x_2$')
ax.set_zlabel('$x_3$')
ax.set_title('Decision Boundary for 3-Variable OR Operation with
```



```
Threshold 1')  
plt.show()
```


Decision Boundary for 3-Variable OR Operation with Threshold 1



Implement X_1 and X_2 with 2-variables and threshold 1.

```
x1 = np.linspace(0, 1, 100)  
x2 = np.linspace(0, 1, 100)  
X1, X2 = np.meshgrid(x1, x2)  
weighted_sum = X1 - X2  
decision_boundary = weighted_sum >= 1  
plt.figure(figsize=(8, 6))
```

```
plt.contourf(X1, X2, decision_boundary, alpha=0.5, cmap='gray',
levels=[0, 1])
plt.plot(X1[0, :], X1[0, :] - 1, 'r--', label='Decision Boundary:  $X_1 - X_2 = 1$ ')
plt.xlabel('$X_1$')
plt.ylabel('$X_2$')
plt.title('Decision Boundary for  $X_1$  AND NOT  $X_2$  with Threshold 1')
plt.xlim(0, 1)
plt.ylim(0, 1)
plt.axhline(0, color='black', linewidth=0.5)
plt.axvline(0, color='black', linewidth=0.5)
plt.grid(True)
plt.legend()
plt.show()
```




Implement NOR with 2-variables and threshold 0.

```
x1 = np.linspace(0, 1, 100)
x2 = np.linspace(0, 1, 100)
```

```

X1, X2 = np.meshgrid(x1, x2)
weighted_sum = -X1 - X2
decision_boundary = weighted_sum >= 0
plt.figure(figsize=(8, 6))
plt.contourf(X1, X2, decision_boundary, alpha=0.5, cmap='gray',
levels=[0, 1])
plt.plot(X1[0, :], -X1[0, :] + 0, 'r--', label='Decision Boundary: $-x_1 - x_2 = 0$')
plt.xlabel('$x_1$')
plt.ylabel('$x_2$')
plt.title('Decision Boundary for 2-Variable NOR Operation')
plt.xlim(0, 1)
plt.ylim(0, 1)
plt.axhline(0, color='black', linewidth=0.5)
plt.axvline(0, color='black', linewidth=0.5)
plt.grid(True)
plt.legend()
plt.show()

```




Implement NOT with 2-variables and threshold 0.

```


x1 = np.linspace(0, 1, 100)
x2 = np.linspace(0, 1, 100)
X1, X2 = np.meshgrid(x1, x2)
weighted_sum = -X1
decision_boundary = weighted_sum >= 0
plt.figure(figsize=(8, 6))
plt.contourf(X1, X2, decision_boundary, alpha=0.5, cmap='gray',
levels=[0, 1])
plt.plot(X1[0, :], -X1[0, :] + 0, 'r--', label='Decision Boundary: $-x_1 = 0$')
plt.xlabel('$x_1$')
plt.ylabel('$x_2$')
plt.title('Decision Boundary for 1-Variable NOT Operation')
plt.xlim(0, 1)
plt.ylim(0, 1)
plt.axhline(0, color='black', linewidth=0.5)
plt.axvline(0, color='black', linewidth=0.5)
plt.grid(True)
plt.legend()
plt.show()

```




Implement OR with 2-variables and threshold 1.

```
def or_gate(x1, x2):  
    return 1 if (x1 + x2) >= 1 else 0  
  
# Generate input values for x1 and x2  
input_combinations = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])  
  
# Calculate the output for each combination of x1 and x2  
outputs = np.array([or_gate(x1, x2) for x1, x2 in input_combinations])  
  
# Display the results  
for (x1, x2), output in zip(input_combinations, outputs):  
    print(f"OR({x1}, {x2}) = {output}")  
  
# Plotting the decision boundary  
# Create a grid for x1 and x2  
x1 = np.linspace(0, 1, 100)  
x2 = np.linspace(0, 1, 100)  
X1, X2 = np.meshgrid(x1, x2)  
  
# Calculate the decision boundary  
Z = (X1 + X2) >= 1  
  
# Plotting  
plt.figure(figsize=(8, 6))  
plt.contourf(X1, X2, Z, alpha=0.5, cmap='coolwarm')  
plt.colorbar(label='Output (1 or 0)')  
plt.scatter(input_combinations[:, 0], input_combinations[:, 1],  
            color='black', label='Input Points')  
  
# Set labels and title  
plt.xlabel('$x_1$')  
plt.ylabel('$x_2$')  
plt.title('2-Variable OR Operation with Threshold 1')  
plt.xlim(0, 1)  
plt.ylim(0, 1)  
plt.axhline(0, color='black', lw=0.5, ls='--')  
plt.axvline(0, color='black', lw=0.5, ls='--')  
plt.legend()  
plt.grid()  
  
# Show the plot  
plt.show()  
  
OR(0, 0) = 0  
OR(0, 1) = 1  
OR(1, 0) = 1  
OR(1, 1) = 1
```



Implement AND with 2-variables and threshold 2.


```
x1 = np.linspace(0, 2, 100)
x2 = 2 - x1
plt.plot(x1, x2, label='Decision Boundary:  $x_1 + x_2 = 2$ ',
color='blue')
plt.scatter([0, 0, 1, 1], [0, 1, 0, 1], color='red')
plt.text(0, 0, 'AND(0,0) = 0', fontsize=10, ha='right')
plt.text(1, 0, 'AND(1,0) = 0', fontsize=10, ha='right')
plt.text(0, 1, 'AND(0,1) = 0', fontsize=10, ha='right')
plt.text(1, 1, 'AND(1,1) = 1', fontsize=10, ha='right')
plt.xlim(0, 2)
plt.ylim(0, 2)
plt.xlabel('$x_1$')
plt.ylabel('$x_2$')
plt.title('Decision Boundary for 2D AND Operation with Threshold 2')
plt.legend()
plt.grid(True)
plt.show()
```



Implement OR with 3-variables and threshold 1.

```
x = np.linspace(0, 1, 100)
X1, X2 = np.meshgrid(x, x)
fig = plt.figure(figsize=(12, 8))
ax = fig.add_subplot(111, projection='3d')
x3_values = [0, 0.5, 1]
for x3 in x3_values:
    Z = (X1 + X2) >= (1 - x3)
    ax.plot_surface(X1, X2, x3 * np.ones_like(X1),
                    facecolors=plt.cm.gray(Z), alpha=0.5,
                    edgecolor='none')
ax.set_xlabel('$x_1$')
ax.set_ylabel('$x_2$')
ax.set_zlabel('$x_3$')
ax.set_title('Decision Boundary for 3-Variable OR Operation with
Threshold 1')
plt.show()
```

Decision Boundary for 3-Variable OR Operation with Threshold 1




Assignment 5

Q1. Implement the OR and AND problem using perceptron learning model.

Algorithm:

Preceptron Network



- transfer function used in the output neuron of the perceptron

$$\bullet f(\text{Sum}_k) = \begin{cases} 1 & \text{if } \text{Sum}_k > \theta_k \\ 0 & \text{if } -\theta_k \leq \text{Sum}_k \leq \theta_k \\ -1 & \text{if } \text{Sum}_k < -\theta_k \end{cases}$$

- $\theta_k = 0.5$, $W_1=0$, $W_2=1$, $n=1$
- Then the weights will be updated

Implement AND and OR.

In each iteration outputs updated weights and t-y and output.

Q2. Implement the same with vector notation.

Where $y=g(W^T X)$ and $W^T = W^T(\text{old}) + \alpha(t-y).X^T$

Q3. Download the Fashion MNIST dataset from

<https://github.com/zalandoresearch/fashion-mnist> and develop an image classification

Q1. Implement the OR and AND problem using perceptron learning model.

```
print("OR PROBLEM")
def g(x):
    if x > 0 :
        return 1
    else : return 0
inputs = [[0,0], [0,1], [1,0], [1,1]]
outputs = [0, 1, 1, 1]
weights = [0.5, 0, 1]
n = 0.25
old_weights = []
iterations = 0
while old_weights!=weights:
    iterations+=1
    old_weights = weights.copy()
    for i in range(4):
        inp, op = inputs[i], outputs[i]
        x, y = inp[0], inp[1]
        val = g(weights[0] + weights[1]*x + weights[2]*y)
        if (val!=op):
            weights[0] = weights[0] + n*(op - val)*1
            weights[1] = weights[1] + n*(op - val)*x
            weights[2] = weights[2] + n*(op - val)*y
    print(weights, '\t', val)
    print(old_weights, '\t', weights, '\n')
print("\nNumber of Iterations are : ", iterations)
```

OR PROBLEM

[0.25, 0.0, 1.0]	1
[0.25, 0.0, 1.0]	1
[0.25, 0.0, 1.0]	1
[0.25, 0.0, 1.0]	1
[0.5, 0, 1]	[0.25, 0.0, 1.0]

[0.0, 0.0, 1.0]	1
[0.0, 0.0, 1.0]	1
[0.25, 0.25, 1.0]	0
[0.25, 0.25, 1.0]	1
[0.25, 0.0, 1.0]	[0.25, 0.25, 1.0]

[0.0, 0.25, 1.0]	1
[0.0, 0.25, 1.0]	1
[0.0, 0.25, 1.0]	1
[0.0, 0.25, 1.0]	1
[0.25, 0.25, 1.0]	[0.0, 0.25, 1.0]

[0.0, 0.25, 1.0]	0
[0.0, 0.25, 1.0]	1

[0.0, 0.25, 1.0]	1
[0.0, 0.25, 1.0]	1
[0.0, 0.25, 1.0]	[0.0, 0.25, 1.0]

Number of Iterations are : 4

```
print("AND PROBLEM")
inputs = [[0,0], [0,1], [1,0], [1,1]]
outputs = [0, 0, 0, 1]
weights = [0.5, 0, 1]
n = 0.25
old_weights = []
iterations = 0
while old_weights!=weights:
    iterations+=1
    old_weights = weights.copy()
    for i in range(4):
        inp, op = inputs[i], outputs[i]
        x, y = inp[0], inp[1]
        val = g(weights[0] + weights[1]*x + weights[2]*y)
        if (val!=op):
            weights[0] = weights[0] + n*(op - val)*1
            weights[1] = weights[1] + n*(op - val)*x
            weights[2] = weights[2] + n*(op - val)*y
    print(weights, '\t', val)
    print(old_weights, '\t', weights, '\n')
print("\nNumber of Iterations are : ", iterations)
```

AND PROBLEM

[0.25, 0.0, 1.0]	1
[0.0, 0.0, 0.75]	1
[0.0, 0.0, 0.75]	0
[0.0, 0.0, 0.75]	1
[0.5, 0, 1]	[0.0, 0.0, 0.75]

[0.0, 0.0, 0.75]	0
[-0.25, 0.0, 0.5]	1
[-0.25, 0.0, 0.5]	0
[-0.25, 0.0, 0.5]	1
[0.0, 0.0, 0.75]	[-0.25, 0.0, 0.5]

[-0.25, 0.0, 0.5]	0
[-0.5, 0.0, 0.25]	1
[-0.5, 0.0, 0.25]	0
[-0.25, 0.25, 0.5]	0
[-0.25, 0.0, 0.5]	[-0.25, 0.25, 0.5]

[-0.25, 0.25, 0.5]	0
[-0.5, 0.25, 0.25]	1

```

[-0.5, 0.25, 0.25]    0
[-0.25, 0.5, 0.5]    0
[-0.25, 0.25, 0.5]    [-0.25, 0.5, 0.5]

[-0.25, 0.5, 0.5]    0
[-0.5, 0.5, 0.25]    1
[-0.5, 0.5, 0.25]    0
[-0.5, 0.5, 0.25]    1
[-0.25, 0.5, 0.5]    [-0.5, 0.5, 0.25]

[-0.5, 0.5, 0.25]    0
[-0.5, 0.5, 0.25]    0
[-0.5, 0.5, 0.25]    0
[-0.5, 0.5, 0.25]    1
[-0.5, 0.5, 0.25]    [-0.5, 0.5, 0.25]

```

Number of Iterations are : 6

Q2. Implement same with vector notation where: $y=g(WTX)$ and $WT=WT(old)+\alpha(t-y).XT$

```

import numpy as np
weights = np.array(((0.5, -0.5), (-0.5, 0), (0.4, -0.2), (-0.1, -
0.1)))
X = np.array(((1, ), (2, ), (1, ), (3, )))
output = (((1, ), (0, )))
n = 0.3
old_weights = np.array(())
activate = np.vectorize(g)
iterations = 1
print(weights.T, '\n')
while old_weights.all() != weights.all():
    iterations += 1
    old_weights = weights.copy()
    y = activate(np.dot(weights.T, X))
    weights = weights + (n * (output - y) * X.T).T
    print(weights.T, '\n')
print("Number of iterations : ", iterations)

[[ 0.5 -0.5  0.4 -0.1]
 [-0.5  0.  -0.2 -0.1]]

[[ 0.8  0.1  0.7  0.8]
 [-0.5  0.  -0.2 -0.1]]

Number of iterations : 2

```

Q3. Download the Fashion MNIST dataset from <https://github.com/zalandoresearch/fashion-mnist> and develop an image classification

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import layers, datasets, models
(train_images, train_labels), (test_images, test_labels) =
datasets.fashion_mnist.load_data()
train_images = train_images / 255.0
test_images = test_images / 255.0
train_images = train_images.reshape((train_images.shape[0], 28, 28,
1))
test_images = test_images.reshape((test_images.shape[0], 28, 28, 1))
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28,
1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
model.fit(train_images, train_labels, epochs=5,
          validation_data=(test_images, test_labels))
history=model.fit(train_images, train_labels, epochs=5,
                  validation_data=(test_images, test_labels))
plt.plot((history.history['accuracy']))
plt.plot((history.history['val_accuracy']))
plt.plot((history.history['loss']))
plt.plot((history.history['val_loss']))
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['accuracy', 'val_accuracy', 'loss', 'val_loss'], loc='upper
left')
plt.show()
test_loss, test_accuracy = model.evaluate(test_images, test_labels)
print(f'\nTest accuracy: {test_accuracy:.4f}')
model.summary()

```

C:\Users\Asus\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

    super().__init__(activity_regularizer=activity_regularizer,
**kwargs)

Epoch 1/5
1875/1875 _____ 61s 29ms/step - accuracy: 0.7458 -
loss: 0.6877 - val_accuracy: 0.8545 - val_loss: 0.3944
Epoch 2/5
1875/1875 _____ 46s 25ms/step - accuracy: 0.8770 -
loss: 0.3343 - val_accuracy: 0.8769 - val_loss: 0.3291
Epoch 3/5
1875/1875 _____ 48s 26ms/step - accuracy: 0.8957 -
loss: 0.2811 - val_accuracy: 0.8922 - val_loss: 0.2931
Epoch 4/5
1875/1875 _____ 47s 25ms/step - accuracy: 0.9088 -
loss: 0.2489 - val_accuracy: 0.8914 - val_loss: 0.2985
Epoch 5/5
1875/1875 _____ 43s 23ms/step - accuracy: 0.9176 -
loss: 0.2213 - val_accuracy: 0.8939 - val_loss: 0.2996
Epoch 1/5
1875/1875 _____ 13s 7ms/step - accuracy: 0.9258 - loss:
0.2017 - val_accuracy: 0.8983 - val_loss: 0.2870
Epoch 2/5
1875/1875 _____ 20s 6ms/step - accuracy: 0.9313 - loss:
0.1847 - val_accuracy: 0.9035 - val_loss: 0.2891
Epoch 3/5
1875/1875 _____ 12s 6ms/step - accuracy: 0.9393 - loss:
0.1665 - val_accuracy: 0.9055 - val_loss: 0.2755
Epoch 4/5
1875/1875 _____ 12s 7ms/step - accuracy: 0.9438 - loss:
0.1477 - val_accuracy: 0.9062 - val_loss: 0.2757
Epoch 5/5
336/1875 _____ 10s 7ms/step - accuracy: 0.9496 - loss:
0.1357

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
predictions = model.predict(test_images)
predicted_labels = np.argmax(predictions, axis=1)
accuracy = accuracy_score(test_labels, predicted_labels)
print(f'Test accuracy: {accuracy:.4f}')
print(classification_report(test_labels, predicted_labels))
conf_matrix = confusion_matrix(test_labels, predicted_labels)
print('Confusion Matrix:')
print(conf_matrix)

class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
               'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
plt.figure(figsize=(10,10))
for i in range(25):

```

```

plt.subplot(5,5,i+1)
plt.xticks([])
plt.yticks([])
plt.grid(False)
plt.imshow(train_images[i], cmap=plt.cm.binary)
plt.xlabel(class_names[train_labels[i]])
plt.show()

def plot_image(i, predictions_array, true_label, img):
    true_label, img = true_label[i], img[i]
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])
    plt.imshow(img, cmap=plt.cm.binary)
    predicted_label = np.argmax(predictions_array)
    if predicted_label == true_label:
        color = 'blue'
    else:
        color = 'red'
    plt.xlabel("{} {:2.0f}% ({})" .format(class_names[predicted_label],
                                         100*np.max(predictions_array),
                                         class_names[true_label]),
              color=color)

def plot_value_array(i, predictions_array, true_label):
    true_label = true_label[i]
    plt.grid(False)
    plt.xticks(range(10))
    plt.yticks([])
    thisplot = plt.bar(range(10), predictions_array, color="#777777")
    plt.ylim([0, 1])
    predicted_label = np.argmax(predictions_array)

    thisplot[predicted_label].set_color('red')
    thisplot[true_label].set_color('blue')

num_rows = 5
num_cols = 3
num_images = num_rows*num_cols
plt.figure(figsize=(2*2*num_cols, 2*2*num_rows))
for i in range(num_images):
    plt.subplot(num_rows, 2*num_cols, 2*i+1)
    plot_image(i, predictions[i], test_labels, test_images)
    plt.subplot(num_rows, 2*num_cols, 2*i+2)
    plot_value_array(i, predictions[i], test_labels)
plt.tight_layout()
plt.show()

```


Assignment 6

Classification of brain MRI using hyper column technique with convolutional neural network and feature selection method

1. The dataset consists of open-access brain tumor MRI containing two classes of the tumor and normal. The dataset consists of 155 and 98 tumor and normal brain MRI, respectively. The dataset is heterogeneous MR images collected from 253 patients.

The images in the dataset has not high resolution and each image has a different resolution. The image format is JPEG.

2. Print the Normal class sample and tumor class sample.

3. Do as follows given in the model below, and compare the results.




Fig. 8. The overall block diagram of the proposed model.

```

import os
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Dropout
from tensorflow.keras.applications import VGG16

dataset_dir = 'images'
img_size = (180, 180)
batch_size = 32
tumor_images = []
no_tumor_images = []

for filename in os.listdir(dataset_dir):
    if filename.lower().startswith('y'):
        tumor_images.append(os.path.join(dataset_dir, filename))
    elif filename.lower().startswith('n'):
        no_tumor_images.append(os.path.join(dataset_dir, filename))

train_tumor_images, val_tumor_images =
tumor_images[:int(0.7*len(tumor_images))],
tumor_images[int(0.7*len(tumor_images)):]
train_no_tumor_images, val_no_tumor_images =
no_tumor_images[:int(0.7*len(no_tumor_images))],
no_tumor_images[int(0.7*len(no_tumor_images)):]

def load_image(file):
    img = tf.io.read_file(file)
    img = tf.image.decode_image(img, channels=3)
    img = tf.image.resize(img, img_size)
    img = img / 255.0
    return img

train_images = [load_image(file) for file in train_tumor_images +
train_no_tumor_images]
train_labels = [1] * len(train_tumor_images) + [0] *
len(train_no_tumor_images)
validation_images = [load_image(file) for file in val_tumor_images +
val_no_tumor_images]
validation_labels = [1] * len(val_tumor_images) + [0] *
len(val_no_tumor_images)

train_datagen = ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,

```

```

        horizontal_flip=True,
        fill_mode='nearest'
    )

validation_datagen = ImageDataGenerator()

train_dataset = train_datagen.flow(
    np.array(train_images),
    np.array(train_labels),
    batch_size=batch_size
)

validation_dataset = validation_datagen.flow(
    np.array(validation_images),
    np.array(validation_labels),
    batch_size=batch_size
)

# CNN
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(180, 180, 3)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.2),
    Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
history = model.fit(train_dataset, epochs=10,
validation_data=validation_dataset)

```

C:\Users\pranj\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

    super().__init__(activity_regularizer=activity_regularizer,
**kwargs)

```

Epoch 1/10

C:\Users\pranj\anaconda3\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:122: UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`,

```
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
```

```
self._warn_if_super_not_called()
```

```
5/5 _____ 5s 526ms/step - accuracy: 0.7473 - loss: 0.6608 - val_accuracy: 0.7581 - val_loss: 0.5011
```

```
Epoch 2/10
```

```
5/5 _____ 3s 478ms/step - accuracy: 0.7929 - loss: 0.5268 - val_accuracy: 0.7581 - val_loss: 0.4939
```

```
Epoch 3/10
```

```
5/5 _____ 3s 600ms/step - accuracy: 0.7586 - loss: 0.5063 - val_accuracy: 0.7581 - val_loss: 0.4909
```

```
Epoch 4/10
```

```
5/5 _____ 4s 541ms/step - accuracy: 0.7309 - loss: 0.5097 - val_accuracy: 0.7581 - val_loss: 0.5064
```

```
Epoch 5/10
```

```
5/5 _____ 3s 523ms/step - accuracy: 0.8106 - loss: 0.4416 - val_accuracy: 0.7581 - val_loss: 0.4950
```

```
Epoch 6/10
```

```
5/5 _____ 3s 494ms/step - accuracy: 0.7642 - loss: 0.4777 - val_accuracy: 0.7581 - val_loss: 0.5321
```

```
Epoch 7/10
```

```
5/5 _____ 3s 499ms/step - accuracy: 0.7930 - loss: 0.5001 - val_accuracy: 0.7581 - val_loss: 0.4818
```

```
Epoch 8/10
```

```
5/5 _____ 3s 554ms/step - accuracy: 0.7468 - loss: 0.5038 - val_accuracy: 0.7581 - val_loss: 0.4827
```

```
Epoch 9/10
```

```
5/5 _____ 4s 635ms/step - accuracy: 0.7787 - loss: 0.4666 - val_accuracy: 0.7903 - val_loss: 0.4593
```


```
Epoch 10/10
```

```
5/5 _____ 4s 531ms/step - accuracy: 0.7433 - loss: 0.5408 - val_accuracy: 0.8387 - val_loss: 0.4553
```

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.show()
```

```
image_path = os.path.join(os.getcwd(), 'images', 'Y10.jpg')
image_string = tf.io.read_file(image_path)
new_image = tf.image.decode_jpeg(image_string, channels=3)
new_image = tf.image.resize(new_image, img_size)
new_image = new_image / 255.0
new_image = new_image[tf.newaxis, ...]
```

```
prediction = model.predict(new_image)
print('Prediction:', prediction)
```



1/1 ————— 0s 80ms/step

Prediction: [[0.9548318]]

dataset_dir = 'images'

img_size = (180, 180)

batch_size = 32

tumor_images = []

no_tumor_images = []

```
for filename in os.listdir(dataset_dir):
    if filename.lower().startswith('y'):
        tumor_images.append(os.path.join(dataset_dir, filename))
    elif filename.lower().startswith('n'):
        no_tumor_images.append(os.path.join(dataset_dir, filename))
```

train_tumor_images = tumor_images[:int(0.8 * len(tumor_images))]

val_tumor_images = tumor_images[int(0.8 * len(tumor_images)):]

train_no_tumor_images = no_tumor_images[:int(0.8 * len(no_tumor_images))]

val_no_tumor_images = no_tumor_images[int(0.8 * len(no_tumor_images)):]

```
train_datagen = ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.2,
```

```

        height_shift_range=0.2,
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True,
        fill_mode='nearest'
    )

    validation_datagen = ImageDataGenerator()

    train_dataset = train_datagen.flow(
        np.array(train_images),
        np.array(train_labels),
        batch_size=batch_size
    )

    validation_dataset = validation_datagen.flow(
        np.array(validation_images),
        np.array(validation_labels),
        batch_size=batch_size
    )

```

AlexNet Model

```

alexnet_model = Sequential([
    Conv2D(96, (11, 11), strides=4, activation='relu',
    input_shape=(180, 180, 3)),
    MaxPooling2D((3, 3), strides=2),
    Conv2D(256, (5, 5), padding='same', activation='relu'),
    MaxPooling2D((3, 3), strides=2),
    Conv2D(384, (3, 3), padding='same', activation='relu'),
    Conv2D(384, (3, 3), padding='same', activation='relu'),
    Conv2D(256, (3, 3), padding='same', activation='relu'),
    MaxPooling2D((3, 3), strides=2),
    Flatten(),
    Dense(4096, activation='relu'),
    Dropout(0.5),
    Dense(4096, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])

alexnet_model.compile(optimizer='adam', loss='binary_crossentropy',
    metrics=['accuracy'])
alexnet_history = alexnet_model.fit(train_dataset, epochs=10,
    validation_data=validation_dataset)

```

C:\Users\pranj\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in

the model instead.

```
super().__init__(activity_regularizer=activity_regularizer,  
**kwargs)
```

Epoch 1/10

5/5 _____ 8s 1s/step - accuracy: 0.5921 - loss: 0.7885
- val_accuracy: 0.7581 - val_loss: 0.8558

Epoch 2/10

5/5 _____ 6s 984ms/step - accuracy: 0.7456 - loss:
0.7222 - val_accuracy: 0.7581 - val_loss: 0.6105

Epoch 3/10

5/5 _____ 8s 1s/step - accuracy: 0.7683 - loss: 0.5766
- val_accuracy: 0.7581 - val_loss: 0.5527

Epoch 4/10

5/5 _____ 7s 1s/step - accuracy: 0.7480 - loss: 0.5603
- val_accuracy: 0.7581 - val_loss: 0.5516

Epoch 5/10

5/5 _____ 7s 1s/step - accuracy: 0.7668 - loss: 0.5449
- val_accuracy: 0.7581 - val_loss: 0.5616

Epoch 6/10

5/5 _____ 6s 1s/step - accuracy: 0.7586 - loss: 0.5425
- val_accuracy: 0.7581 - val_loss: 0.5774

Epoch 7/10

5/5 _____ 6s 1s/step - accuracy: 0.7454 - loss: 0.5838
- val_accuracy: 0.7581 - val_loss: 0.5649

Epoch 8/10

5/5 _____ 8s 1s/step - accuracy: 0.7286 - loss: 0.5866
- val_accuracy: 0.7581 - val_loss: 0.5533

Epoch 9/10

5/5 _____ 8s 1s/step - accuracy: 0.7835 - loss: 0.5300
- val_accuracy: 0.7581 - val_loss: 0.5579

Epoch 10/10

5/5 _____ 6s 1s/step - accuracy: 0.7938 - loss: 0.5099
- val_accuracy: 0.7581 - val_loss: 0.5530

```
def plot_training_history(history, model_name):  
    plt.figure(figsize=(12, 4))
```

```
    # Plot 1
```

```
    plt.subplot(1, 2, 1)
```

```
    plt.plot(history.history['accuracy'], label='Training Accuracy')
```

```
    plt.plot(history.history['val_accuracy'], label='Validation
```

```
Accuracy')
```

```
    plt.title(f'{model_name} - Accuracy')
```

```
    plt.xlabel('Epoch')
```

```
    plt.ylabel('Accuracy')
```

```
    plt.legend()
```

```
    # Plot 2
```

```
    plt.subplot(1, 2, 2)
```



```

plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title(f'{model_name} - Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.show()

plot_training_history(alexnet_history, 'AlexNet')

```



VGG16

```

vgg16_base = VGG16(weights='imagenet', include_top=False,
input_shape=(180, 180, 3))
for layer in vgg16_base.layers:
    layer.trainable = False

vgg16_model = Sequential([
    vgg16_base,
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])

vgg16_model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
vgg16_history = vgg16_model.fit(train_dataset, epochs=10,
validation_data=validation_dataset)

```

Epoch 1/10

5/5 ————— 25s 5s/step - accuracy: 0.5929 - loss: 2.3551
- val_accuracy: 0.7581 - val_loss: 0.7729


Epoch 2/10


```

5/5 _____ 28s 6s/step - accuracy: 0.6697 - loss: 0.7186
- val_accuracy: 0.7742 - val_loss: 0.4641
Epoch 3/10
5/5 _____ 27s 5s/step - accuracy: 0.7140 - loss: 0.6636
- val_accuracy: 0.7581 - val_loss: 0.4600
Epoch 4/10
5/5 _____ 23s 5s/step - accuracy: 0.7784 - loss: 0.5030
- val_accuracy: 0.7419 - val_loss: 0.5739
Epoch 5/10
5/5 _____ 24s 5s/step - accuracy: 0.7592 - loss: 0.4913
- val_accuracy: 0.7903 - val_loss: 0.4043
Epoch 6/10
5/5 _____ 24s 5s/step - accuracy: 0.7179 - loss: 0.5517
- val_accuracy: 0.7742 - val_loss: 0.3874
Epoch 7/10
5/5 _____ 23s 5s/step - accuracy: 0.7935 - loss: 0.4438
- val_accuracy: 0.7742 - val_loss: 0.3997
Epoch 8/10
5/5 _____ 25s 5s/step - accuracy: 0.7863 - loss: 0.4781
- val_accuracy: 0.9194 - val_loss: 0.4357
Epoch 9/10
5/5 _____ 28s 6s/step - accuracy: 0.7980 - loss: 0.4262
- val_accuracy: 0.8387 - val_loss: 0.3962
Epoch 10/10
5/5 _____ 28s 6s/step - accuracy: 0.8362 - loss: 0.3654
- val_accuracy: 0.8710 - val_loss: 0.3902

plot_training_history(vgg16_history, 'VGG16')

```



Augmentation

```

def display_augmented_images(datagen, image_path):
    if not os.path.exists(image_path):
        print(f"File not found: {image_path}")
        return

```

```

image_string = tf.io.read_file(image_path)
image = tf.image.decode_jpeg(image_string, channels=3)
image = tf.image.resize(image, img_size)
image = image / 255.0
image = image[tf.newaxis, ...]

augmented_images = datagen.flow(image, batch_size=1)










plt.figure(figsize=(12, 12))
for i in range(9):
    plt.subplot(3, 3, i + 1)
    batch = next(augmented_images)
    augmented_image = batch[0]
    plt.imshow(augmented_image)
    plt.axis('off')

plt.show()

image_to_display = 'Y12.jpg'

if image_to_display:
    display_augmented_images(train_datagen, os.path.join(dataset_dir,
image_to_display))
else:
    print("No tumor images found in the dataset.")

```



```
image_path = os.path.join(os.getcwd(), 'images', 'N11.jpg')
image_string = tf.io.read_file(image_path)
new_image = tf.image.decode_jpeg(image_string, channels=3)
new_image = tf.image.resize(new_image, img_size)
new_image = new_image / 255.0
new_image = new_image[tf.newaxis, ...]

vgg16_prediction = vgg16_model.predict(new_image)
alexnet_prediction = alexnet_model.predict(new_image)

print('VGG16 Prediction:', vgg16_prediction)
print('AlexNet Prediction:', alexnet_prediction)
```

```
1/1 _____ 1s 660ms/step
1/1 _____ 0s 136ms/step
VGG16 Prediction: [[0.65607154]]
AlexNet Prediction: [[0.75484836]]
```

Assignment 7

1. Implement Logistic Regression with Autoencoder on MNIST Dataset
2. Add some noise to images and then apply the autoencoder to denoise the dataset.

NAME - VISHAL MARANDI | ROLL - BTECH/10119/21 | DL ASSIGNMENT DATED 16th OCTOBER 2024

Implementing Logistic Regression with Autoencoder on MNIST DATASET

```
!pip install tensorflow
```

```
Requirement already satisfied: tensorflow in
/usr/local/lib/python3.10/dist-packages (2.17.0)
Requirement already satisfied: absl-py>=1.0.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (24.3.25)
Requirement already satisfied: gast!=0.5.0,!=0.5.1,!=0.5.2,>=0.2.1
in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: google-pasta>=0.1.1 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: h5py>=3.10.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (3.11.0)
Requirement already satisfied: libclang>=13.0.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: ml-dtypes<0.5.0,>=0.3.1 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (0.4.1)
Requirement already satisfied: opt-einsum>=2.3.2 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: packaging in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (24.1)
Requirement already satisfied: protobuf!=4.21.0,!=4.21.1,!=4.21.2,!
=4.21.3,!=4.21.4,!=4.21.5,<5.0.0dev,>=3.20.3 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (3.20.3)
Requirement already satisfied: requests<3,>=2.21.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (2.32.3)
Requirement already satisfied: setuptools in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (71.0.4)
Requirement already satisfied: six>=1.12.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: termcolor>=1.1.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (2.5.0)
Requirement already satisfied: typing-extensions>=3.6.6 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (4.12.2)
Requirement already satisfied: wrapt>=1.11.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (1.64.1)
Requirement already satisfied: tensorboard<2.18,>=2.17 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (2.17.0)
Requirement already satisfied: keras>=3.2.0 in
```

/usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem<=0.23.1 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (0.37.1)
Requirement already satisfied: numpy<2.0.0,>=1.23.5 in
/usr/local/lib/python3.10/dist-packages (from tensorflow) (1.26.4)
Requirement already satisfied: wheel<1.0,>=0.23.0 in
/usr/local/lib/python3.10/dist-packages (from astunparse>=1.6.0-
>tensorflow) (0.44.0)
Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-
packages (from keras>=3.2.0->tensorflow) (13.9.2)
Requirement already satisfied: namex in
/usr/local/lib/python3.10/dist-packages (from keras>=3.2.0-
>tensorflow) (0.0.8)
Requirement already satisfied: optree in
/usr/local/lib/python3.10/dist-packages (from keras>=3.2.0-
>tensorflow) (0.13.0)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
>tensorflow) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
>tensorflow) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
>tensorflow) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
>tensorflow) (2024.8.30)
Requirement already satisfied: markdown>=2.6.8 in
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17-
>tensorflow) (3.7)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0
in /usr/local/lib/python3.10/dist-packages (from
tensorboard<2.18,>=2.17->tensorflow) (0.7.2)
Requirement already satisfied: werkzeug>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17-
>tensorflow) (3.0.4)
Requirement already satisfied: MarkupSafe>=2.1.1 in
/usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1-
>tensorboard<2.18,>=2.17->tensorflow) (3.0.1)
Requirement already satisfied: markdown-it-py>=2.2.0 in
/usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0-
>tensorflow) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0-
>tensorflow) (2.18.0)
Requirement already satisfied: mdurl~=0.1 in
/usr/local/lib/python3.10/dist-packages (from markdown-it-py>=2.2.0-
>rich->keras>=3.2.0->tensorflow) (0.1.2)

```

#Importing Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tensorflow.keras.layers import Input,Dense,Flatten,Reshape
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

#Loading the data
(x_train,y_train),(x_test,y_test)=mnist.load_data()

#Preprocessing the data
x_train=x_train.astype('float32')/255
x_test=x_test.astype('float32')/255

#Flattening
x_train_flatten=x_train.reshape(-1,28*28)
x_test_flatten=x_test.reshape(-1,28*28)

# 3. Build Autoencoder Model
input_img = Input(shape=(28 * 28,))
encoded = Dense(128, activation='relu')(input_img) # Change this to
128 features
decoded = Dense(28 * 28, activation='sigmoid')(encoded)
# Autoencoder model
autoencoder = Model(input_img, decoded)
# Encoder model to extract encoded features
encoder = Model(input_img, encoded)

# 4. Compile the Model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
# 5. Train the Autoencoder
autoencoder.fit(x_train_flatten, x_train_flatten,
epochs=20,batch_size=256,shuffle=True,validation_data=(x_test_flatten,
x_test_flatten))
history=autoencoder.fit(x_train_flatten, x_train_flatten,
epochs=5,batch_size=256,shuffle=True,validation_data=(x_test_flatten,
x_test_flatten))

plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

# 6. Use Encoded Features for Logistic Regression

```



```

x_train_encoded = encoder.predict(x_train_flatten)
x_test_encoded = encoder.predict(x_test_flatten)
logistic = LogisticRegression(max_iter=1000)
logistic.fit(x_train_encoded, y_train)
# Predictions
y_pred = logistic.predict(x_test_encoded)
# 7. Evaluate and Analyze
accuracy = accuracy_score(y_test, y_pred)
print(f'Logistic Regression accuracy on encoded features: {accuracy *
100:.2f}%')

```

```

Epoch 1/20
235/235 ————— 6s 19ms/step - loss: 0.0656 - val_loss:
0.0654
Epoch 2/20
235/235 ————— 5s 19ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 3/20
235/235 ————— 7s 30ms/step - loss: 0.0655 - val_loss:
0.0654
Epoch 4/20
235/235 ————— 7s 29ms/step - loss: 0.0655 - val_loss:
0.0654
Epoch 5/20
235/235 ————— 10s 29ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 6/20
235/235 ————— 5s 23ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 7/20
235/235 ————— 11s 25ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 8/20
235/235 ————— 5s 23ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 9/20
235/235 ————— 8s 35ms/step - loss: 0.0655 - val_loss:
0.0654
Epoch 10/20
235/235 ————— 10s 32ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 11/20
235/235 ————— 7s 32ms/step - loss: 0.0656 - val_loss:
0.0654
Epoch 12/20
235/235 ————— 10s 31ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 13/20
235/235 ————— 8s 23ms/step - loss: 0.0653 - val_loss:
0.0654

```

```
Epoch 14/20
235/235 _____ 7s 29ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 15/20
235/235 _____ 4s 15ms/step - loss: 0.0654 - val_loss:
0.0653
Epoch 16/20
235/235 _____ 4s 17ms/step - loss: 0.0654 - val_loss:
0.0653
Epoch 17/20
235/235 _____ 9s 38ms/step - loss: 0.0655 - val_loss:
0.0653
Epoch 18/20
235/235 _____ 7s 31ms/step - loss: 0.0654 - val_loss:
0.0654
Epoch 19/20
235/235 _____ 6s 15ms/step - loss: 0.0653 - val_loss:
0.0653
Epoch 20/20
235/235 _____ 4s 15ms/step - loss: 0.0653 - val_loss:
0.0653
Epoch 1/5
235/235 _____ 5s 23ms/step - loss: 0.0654 - val_loss:
0.0653
Epoch 2/5
235/235 _____ 8s 15ms/step - loss: 0.0655 - val_loss:
0.0653
Epoch 3/5
235/235 _____ 7s 24ms/step - loss: 0.0653 - val_loss:
0.0653
Epoch 4/5
235/235 _____ 10s 24ms/step - loss: 0.0654 - val_loss:
0.0653
Epoch 5/5
235/235 _____ 7s 27ms/step - loss: 0.0656 - val_loss:
0.0654
```



1875/1875 ————— 4s 2ms/step

313/313 ————— 1s 2ms/step

Logistic Regression accuracy on encoded features: 92.24%

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
decoded_imgs=autoencoder.predict(x_test_flatten)
```

```
x_test_resaped=x_test_flatten.reshape(-1,28,28)
```


```
decoded_imgs=decoded_imgs.reshape(-1,28,28)
```

```
decoded_imgs_resaped=decoded_imgs.reshape(-1,28,28)
```

```
n=10
```

```
plt.figure(figsize=(20,4))
for i in range(n):
    ax=plt.subplot(2,n,i+1)
    plt.imshow(x_test_resaped[i],cmap='gray')
    plt.title("Original Image")
    plt.axis('off')
    ax=plt.subplot(2,n,i+1+n)
    plt.imshow(decoded_imgs_resaped[i],cmap='gray')
    plt.title("Reconstructed Image")
    plt.axis('off')
plt.show()
```

313/313 ————— 1s 3ms/step



1. Import Libraries

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
```

2. Load and Preprocess the Dataset

```
(x_train, _), (x_test, _) = mnist.load_data()
```

Normalize the images to [0, 1]

```
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
```

Flatten the images (to pass through the fully connected autoencoder)

```
x_train_flatten = x_train.reshape(-1, 28 * 28)
x_test_flatten = x_test.reshape(-1, 28 * 28)
```

3. Add noise to the dataset

```
noise_factor = 0.75 # Change this to increase or decrease noise level
x_train_noisy = x_train_flatten + noise_factor *
np.random.normal(loc=0.0, scale=1.0, size=x_train_flatten.shape)
x_test_noisy = x_test_flatten + noise_factor *
np.random.normal(loc=0.0, scale=1.0, size=x_test_flatten.shape)
```

Clip the values to be between 0 and 1 (valid pixel range)

```

x_train_noisy = np.clip(x_train_noisy, 0., 1.)
x_test_noisy = np.clip(x_test_noisy, 0., 1.)

# 4. Build Autoencoder Model
input_img = Input(shape=(28 * 28,))
encoded = Dense(128, activation='relu')(input_img)
decoded = Dense(28 * 28, activation='sigmoid')(encoded)

# Autoencoder model
autoencoder = Model(input_img, decoded)

# 5. Compile the Model
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# 6. Train the Autoencoder
history = autoencoder.fit(x_train_noisy, x_train_flatten,
                          epochs=5,
                          batch_size=256,
                          shuffle=True,
                          validation_data=(x_test_noisy,
x_test_flatten))

# 7. Predict Denoised Images
denoised_imgs = autoencoder.predict(x_test_noisy)

# Reshape the noisy and denoised images back to 28x28 for
visualization
x_test_noisy_resaped = x_test_noisy.reshape(-1, 28, 28)
denoised_imgs_resaped = denoised_imgs.reshape(-1, 28, 28)

# 8. Visualize Original Noisy and Denoised Images
n = 10 # number of images to display
plt.figure(figsize=(20, 6))
for i in range(n):
    # Display noisy images
    ax = plt.subplot(3, n, i + 1)
    plt.imshow(x_test_noisy_resaped[i], cmap='gray')
    plt.title("Noisy")
    plt.axis('off')


    # Display denoised images
    ax = plt.subplot(3, n, i + 1 + n)
    plt.imshow(denoised_imgs_resaped[i], cmap='gray')
    plt.title("Denoised")
    plt.axis('off')

    # Display original images for comparison
    ax = plt.subplot(3, n, i + 1 + 2 * n)
    plt.imshow(x_test[i], cmap='gray')
    plt.title("Original")

```

```
plt.axis('off')  
plt.show()
```

Epoch 1/5
235/235 ————— 5s 16ms/step - loss: 0.3215 - val_loss: 0.1912
Epoch 2/5
235/235 ————— 5s 16ms/step - loss: 0.1845 - val_loss: 0.1669
Epoch 3/5
235/235 ————— 5s 16ms/step - loss: 0.1648 - val_loss: 0.1568
Epoch 4/5
235/235 ————— 5s 15ms/step - loss: 0.1556 - val_loss: 0.1513
Epoch 5/5
235/235 ————— 5s 20ms/step - loss: 0.1504 - val_loss: 0.1479
313/313 ————— 1s 3ms/step



Assignment 8

	Lab Assignment Deep Learning Lab
Q1(a)	You have to collect daily temperature (in degrees) of Mesra at 10 AM for last 60 days. Store data in temp.csv file.
(b)	Normalize the data points. Use Min-Max scaling
©	Create a sequence of data points based on last three days to predict the next day For example normalized_temps is [0.1, 0.2, 0.3, 0.4, 0.5, 0.6], and n_input_days is 3. For i = 0: X becomes [[0.1, 0.2, 0.3]] y becomes [0.4] For i = 1: X becomes [[0.1, 0.2, 0.3], [0.2, 0.3, 0.4]] y becomes [0.4, 0.5] similarly for others
(d)	You have to build LSTM model to predict the temperature of next day based on previous days temperature. Use MSE as the loss function and optimizer 'adam'. LSTMs require 3D input because they process data sequentially over time
(e)	Train your model for 50 epoch.
(f)	Use the model to predict values
	import numpy as np from keras.models import Sequential from keras.layers import LSTM, Dense
Q2	Predict the next number in a sequence using RNN. Use MSE as the loss function and optimizer 'adam'.
Q3	Predict the next character in a sequence using RNN given a string of characters. Use sparse_categorical_crossentropy as the loss function and optimizer 'adam'.

Required Libraries

```
import numpy as np
import pandas as pd
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
```

Load the data (Q1a)

```
data = pd.read_csv('/content/temp.csv')
print("Original Data:")
print(data.head())
```

Original Data:

	date	temperature_2m	relative_humidity_2m	\
0	2024-10-30 00:00:00+00:00	21.776499	92.0	
1	2024-10-30 01:00:00+00:00	21.826500	91.0	
2	2024-10-30 02:00:00+00:00	23.076500	86.0	
3	2024-10-30 03:00:00+00:00	24.726500	81.0	
4	2024-10-30 04:00:00+00:00	26.376500	77.0	

	apparent_temperature	precipitation_probability	precipitation	\
0	25.369894	5.0	0.0	
1	25.363144	15.0	0.0	
2	26.651030	13.0	0.0	
3	28.552795	13.0	0.0	
4	30.727776	28.0	0.0	

	weather_code	pressure_msl	surface_pressure	cloud_cover	\
0	1.0	1010.6	940.59380	53.0	
1	2.0	1011.7	941.62900	52.0	
2	2.0	1012.5	942.65704	71.0	
3	2.0	1013.4	943.86630	47.0	
4	2.0	1013.4	944.23350	58.0	

visibility	wind_speed_120m	wind_gusts_10m	uv_index
------------	-----------------	----------------	----------

0	24140.0	11.246759	8.640000	0.00
1	24140.0	11.609651	9.360000	0.10
2	24140.0	12.722830	10.440001	0.85
3	24140.0	5.815978	12.959999	1.70
4	24140.0	5.052841	11.879999	2.95

Q1(b): Normalize the data using Min-Max scaling

```
scaler = MinMaxScaler()
temperature_data = data['temperature_2m'].values.reshape(-1, 1)
normalized_temps = scaler.fit_transform(temperature_data)
```

```
print("\nNormalized Temperature Data (first 5 values):")
print(normalized_temps[:5])
```

Normalized Temperature Data (first 5 values):

```
[[0.38805478]
 [0.3921345 ]
 [0.49412537]
 [0.62875332]
 [0.76338126]]
```

Q1(c): Create sequences for the last three days (72 hours) to predict next day

```
def create_sequences(data, seq_length):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:(i + seq_length)])
        y.append(data[i + seq_length])
    return np.array(X), np.array(y)
```

```
seq_length = 72 # 3 days of hourly data
X, y = create_sequences(normalized_temps, seq_length)
```

```
print("\nSequence Shape:", X.shape)
print("Target Shape:", y.shape)
```

```
Sequence Shape: (312, 72, 1)
Target Shape: (312, 1)
```

Q1(d): Build LSTM model

```
model = Sequential([
    LSTM(50, activation='relu', input_shape=(seq_length, 1),
        return_sequences=True),
    LSTM(50, activation='relu'),
    Dense(1)
])
```

```
model.compile(optimizer='adam', loss='mse')
```

```
print("\nModel Summary:")
model.summary()
```

Model Summary:

```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/
rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
```

Model: "sequential_1"

Layer (type) Param #	Output Shape
lstm_2 (LSTM) 10,400	(None, 72, 50)
lstm_3 (LSTM) 20,200	(None, 50)
dense_1 (Dense) 51	(None, 1)

Total params: 30,651 (119.73 KB)

Trainable params: 30,651 (119.73 KB)

Non-trainable params: 0 (0.00 B)

```
# Q1(e): Train the model
# Reshape input data for LSTM [samples, timesteps, features]
X = X.reshape((X.shape[0], X.shape[1], 1))
```

```
history = model.fit(X, y,
                    epochs=50,
                    batch_size=32,
                    validation_split=0.2,
                    verbose=1)
```

Epoch 1/50

8/8 ————— 4s 112ms/step - loss: 0.3183 - val_loss:

```
0.1674
Epoch 2/50
8/8 _____ 1s 64ms/step - loss: 0.2261 - val_loss:
0.1040
Epoch 3/50
8/8 _____ 1s 61ms/step - loss: 0.1084 - val_loss:
0.1049
Epoch 4/50
8/8 _____ 1s 82ms/step - loss: 0.0923 - val_loss:
0.0687
Epoch 5/50
8/8 _____ 1s 94ms/step - loss: 0.0891 - val_loss:
0.0681
Epoch 6/50
8/8 _____ 1s 90ms/step - loss: 0.0756 - val_loss:
0.0758
Epoch 7/50
8/8 _____ 1s 91ms/step - loss: 0.0780 - val_loss:
0.0705
Epoch 8/50
8/8 _____ 0s 61ms/step - loss: 0.0666 - val_loss:
0.0687
Epoch 9/50
8/8 _____ 1s 63ms/step - loss: 0.0698 - val_loss:
0.0711
Epoch 10/50
8/8 _____ 1s 60ms/step - loss: 0.0670 - val_loss:
0.0670
Epoch 11/50
8/8 _____ 1s 63ms/step - loss: 0.0677 - val_loss:
0.0697
Epoch 12/50
8/8 _____ 1s 59ms/step - loss: 0.0715 - val_loss:
0.0690
Epoch 13/50
8/8 _____ 1s 63ms/step - loss: 0.0681 - val_loss:
0.0666
Epoch 14/50
8/8 _____ 1s 57ms/step - loss: 0.0687 - val_loss:
0.0693
Epoch 15/50
8/8 _____ 1s 58ms/step - loss: 0.0666 - val_loss:
0.0646
Epoch 16/50
8/8 _____ 0s 59ms/step - loss: 0.0651 - val_loss:
0.0694
Epoch 17/50
8/8 _____ 1s 60ms/step - loss: 0.0648 - val_loss:
0.0627
```

```
Epoch 18/50
8/8 _____ 1s 63ms/step - loss: 0.0592 - val_loss:
0.0675
Epoch 19/50
8/8 _____ 0s 59ms/step - loss: 0.0604 - val_loss:
0.0624
Epoch 20/50
8/8 _____ 1s 62ms/step - loss: 0.0580 - val_loss:
0.0618
Epoch 21/50
8/8 _____ 1s 60ms/step - loss: 0.0550 - val_loss:
0.0626
Epoch 22/50
8/8 _____ 1s 63ms/step - loss: 0.0539 - val_loss:
0.0585
Epoch 23/50
8/8 _____ 1s 62ms/step - loss: 0.0538 - val_loss:
0.0685
Epoch 24/50
8/8 _____ 1s 62ms/step - loss: 0.0536 - val_loss:
0.0598
Epoch 25/50
8/8 _____ 1s 59ms/step - loss: 0.0510 - val_loss:
0.0572
Epoch 26/50
8/8 _____ 1s 95ms/step - loss: 0.0487 - val_loss:
0.0488
Epoch 27/50
8/8 _____ 1s 94ms/step - loss: 0.0533 - val_loss:
0.0488
Epoch 28/50
8/8 _____ 1s 86ms/step - loss: 0.0534 - val_loss:
0.0632
Epoch 29/50
8/8 _____ 1s 59ms/step - loss: 0.0476 - val_loss:
0.0483
Epoch 30/50
8/8 _____ 1s 60ms/step - loss: 0.0471 - val_loss:
0.0758
Epoch 31/50
8/8 _____ 0s 58ms/step - loss: 0.0433 - val_loss:
0.0436
Epoch 32/50
8/8 _____ 0s 60ms/step - loss: 0.0435 - val_loss:
0.0549
Epoch 33/50
8/8 _____ 1s 61ms/step - loss: 0.0322 - val_loss:
0.0475
Epoch 34/50
```

```
8/8 _____ 1s 65ms/step - loss: 0.0327 - val_loss:
0.0402
Epoch 35/50
8/8 _____ 0s 59ms/step - loss: 0.0202 - val_loss:
0.0316
Epoch 36/50
8/8 _____ 1s 60ms/step - loss: 0.0186 - val_loss:
0.0153
Epoch 37/50
8/8 _____ 0s 61ms/step - loss: 0.0173 - val_loss:
0.0209
Epoch 38/50
8/8 _____ 0s 60ms/step - loss: 0.0183 - val_loss:
0.0354
Epoch 39/50
8/8 _____ 1s 61ms/step - loss: 0.0128 - val_loss:
0.0142
Epoch 40/50
8/8 _____ 1s 61ms/step - loss: 0.0102 - val_loss:
0.0102
Epoch 41/50
8/8 _____ 1s 61ms/step - loss: 0.0090 - val_loss:
0.0137
Epoch 42/50
8/8 _____ 1s 62ms/step - loss: 0.0073 - val_loss:
0.0100
Epoch 43/50
8/8 _____ 1s 61ms/step - loss: 0.0068 - val_loss:
0.0089
Epoch 44/50
8/8 _____ 1s 63ms/step - loss: 0.0068 - val_loss:
0.0073
Epoch 45/50
8/8 _____ 1s 82ms/step - loss: 0.0055 - val_loss:
0.0074
Epoch 46/50
8/8 _____ 1s 96ms/step - loss: 0.0050 - val_loss:
0.0072
Epoch 47/50
8/8 _____ 1s 101ms/step - loss: 0.0056 - val_loss:
0.0070
Epoch 48/50
8/8 _____ 1s 63ms/step - loss: 0.0051 - val_loss:
0.0076
Epoch 49/50
8/8 _____ 0s 60ms/step - loss: 0.0043 - val_loss:
0.0049
Epoch 50/50
```

8/8 ————— 1s 65ms/step - loss: 0.0039 - val_loss: 0.0056

```
# Q1(f): Make predictions
# Prepare last sequence for prediction
last_sequence = normalized_temps[-seq_length:]
last_sequence = last_sequence.reshape((1, seq_length, 1))

# Make prediction
predicted_normalized = model.predict(last_sequence)

# Denormalize prediction
predicted_temp = scaler.inverse_transform(predicted_normalized)


print("\nPrediction Results:")
print(f"Predicted next temperature: {predicted_temp[0][0]:.2f}°C")

# Plot training history
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss During Training')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

1/1 ————— 0s 436ms/step

Prediction Results:
Predicted next temperature: 19.46°C



```
import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.preprocessing import MinMaxScaler

# Define the sequence
sequence = np.array([1, 2, 4, 8, 16, 32, 64, 128, 256, 512])

# Normalize the data
scaler = MinMaxScaler()
normalized_sequence = scaler.fit_transform(sequence.reshape(-1, 1))

# Create sequences for training
def create_sequences(data, seq_length):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:(i + seq_length)])
        y.append(data[i + seq_length])
    return np.array(X), np.array(y)

seq_length = 3
X, y = create_sequences(normalized_sequence, seq_length)

# Reshape input to be [samples, time steps, features]
X = X.reshape((X.shape[0], X.shape[1], 1))
```

```

# Define the model
model = Sequential([
    LSTM(50, activation='relu', input_shape=(seq_length, 1),
return_sequences=True),
    LSTM(50, activation='relu'),
    Dense(1)
])

# Compile the model
model.compile(optimizer='adam', loss='mse')

# Train the model
history = model.fit(X, y, epochs=100, batch_size=1, verbose=1)

# Make a prediction
last_sequence = normalized_sequence[-seq_length:]
last_sequence = last_sequence.reshape((1, seq_length, 1))
predicted_normalized = model.predict(last_sequence)

# Denormalize the prediction
predicted = scaler.inverse_transform(predicted_normalized)

print(f"Sequence: {sequence}")
print(f"Predicted next number: {predicted[0][0]:.2f}")

# Plot the training loss
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'])
plt.title('Model Loss During Training')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show()

```

Epoch 1/100

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/
rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

```

7/7 ————— 2s 3ms/step - loss: 0.1505

Epoch 2/100

7/7 ————— 0s 3ms/step - loss: 0.1349

Epoch 3/100

7/7 ————— 0s 4ms/step - loss: 0.3323

Epoch 4/100

7/7 ————— 0s 3ms/step - loss: 0.0996


Epoch 5/100


```
7/7 _____ 0s 3ms/step - loss: 0.1356
Epoch 6/100
7/7 _____ 0s 4ms/step - loss: 0.1225
Epoch 7/100
7/7 _____ 0s 4ms/step - loss: 0.1377
Epoch 8/100
7/7 _____ 0s 3ms/step - loss: 0.0573
Epoch 9/100
7/7 _____ 0s 3ms/step - loss: 0.0725
Epoch 10/100
7/7 _____ 0s 3ms/step - loss: 0.0569
Epoch 11/100
7/7 _____ 0s 3ms/step - loss: 0.0872
Epoch 12/100
7/7 _____ 0s 3ms/step - loss: 0.0462
Epoch 13/100
7/7 _____ 0s 3ms/step - loss: 0.0439
Epoch 14/100
7/7 _____ 0s 4ms/step - loss: 0.0649
Epoch 15/100
7/7 _____ 0s 3ms/step - loss: 0.0899
Epoch 16/100
7/7 _____ 0s 4ms/step - loss: 0.1173
Epoch 17/100
7/7 _____ 0s 4ms/step - loss: 0.1440
Epoch 18/100
7/7 _____ 0s 3ms/step - loss: 0.0646
Epoch 19/100
7/7 _____ 0s 3ms/step - loss: 0.0682
Epoch 20/100
7/7 _____ 0s 3ms/step - loss: 0.1061
Epoch 21/100
7/7 _____ 0s 3ms/step - loss: 0.0370
Epoch 22/100
7/7 _____ 0s 3ms/step - loss: 0.0718
Epoch 23/100
7/7 _____ 0s 3ms/step - loss: 0.0367
Epoch 24/100
7/7 _____ 0s 4ms/step - loss: 0.0450
Epoch 25/100
7/7 _____ 0s 5ms/step - loss: 0.0317
Epoch 26/100
7/7 _____ 0s 4ms/step - loss: 0.0337
Epoch 27/100
7/7 _____ 0s 4ms/step - loss: 0.0856
Epoch 28/100
7/7 _____ 0s 4ms/step - loss: 0.0474
Epoch 29/100
7/7 _____ 0s 4ms/step - loss: 0.0362
```

```
Epoch 30/100
7/7 _____ 0s 3ms/step - loss: 0.0229
Epoch 31/100
7/7 _____ 0s 3ms/step - loss: 0.0292
Epoch 32/100
7/7 _____ 0s 5ms/step - loss: 0.0150
Epoch 33/100
7/7 _____ 0s 3ms/step - loss: 0.0112
Epoch 34/100
7/7 _____ 0s 3ms/step - loss: 0.0073
Epoch 35/100
7/7 _____ 0s 3ms/step - loss: 0.0064
Epoch 36/100
7/7 _____ 0s 3ms/step - loss: 0.0029
Epoch 37/100
7/7 _____ 0s 3ms/step - loss: 0.0022
Epoch 38/100
7/7 _____ 0s 3ms/step - loss: 8.7448e-04
Epoch 39/100
7/7 _____ 0s 3ms/step - loss: 0.0014
Epoch 40/100
7/7 _____ 0s 3ms/step - loss: 9.4513e-04
Epoch 41/100
7/7 _____ 0s 4ms/step - loss: 5.3487e-04
Epoch 42/100
7/7 _____ 0s 4ms/step - loss: 5.6396e-04
Epoch 43/100
7/7 _____ 0s 3ms/step - loss: 9.1996e-04
Epoch 44/100
7/7 _____ 0s 3ms/step - loss: 0.0010
Epoch 45/100
7/7 _____ 0s 4ms/step - loss: 4.5551e-04
Epoch 46/100
7/7 _____ 0s 4ms/step - loss: 0.0011
Epoch 47/100
7/7 _____ 0s 3ms/step - loss: 8.5797e-04
Epoch 48/100
7/7 _____ 0s 3ms/step - loss: 5.5944e-04
Epoch 49/100
7/7 _____ 0s 5ms/step - loss: 3.3262e-04
Epoch 50/100
7/7 _____ 0s 3ms/step - loss: 7.9507e-04
Epoch 51/100
7/7 _____ 0s 3ms/step - loss: 5.7610e-04
Epoch 52/100
7/7 _____ 0s 3ms/step - loss: 6.8836e-04
Epoch 53/100
7/7 _____ 0s 3ms/step - loss: 4.1971e-04
Epoch 54/100
```

```
7/7 _____ 0s 3ms/step - loss: 2.4971e-04
Epoch 55/100
7/7 _____ 0s 4ms/step - loss: 3.3731e-04
Epoch 56/100
7/7 _____ 0s 3ms/step - loss: 6.7571e-04
Epoch 57/100
7/7 _____ 0s 3ms/step - loss: 3.7042e-04
Epoch 58/100
7/7 _____ 0s 4ms/step - loss: 5.4504e-04
Epoch 59/100
7/7 _____ 0s 4ms/step - loss: 4.0497e-04
Epoch 60/100
7/7 _____ 0s 4ms/step - loss: 3.4908e-04
Epoch 61/100
7/7 _____ 0s 4ms/step - loss: 4.2530e-04
Epoch 62/100
7/7 _____ 0s 3ms/step - loss: 2.6994e-04
Epoch 63/100
7/7 _____ 0s 4ms/step - loss: 3.7053e-04
Epoch 64/100
7/7 _____ 0s 4ms/step - loss: 2.6442e-04
Epoch 65/100
7/7 _____ 0s 4ms/step - loss: 3.4208e-04
Epoch 66/100
7/7 _____ 0s 3ms/step - loss: 2.0784e-04
Epoch 67/100
7/7 _____ 0s 4ms/step - loss: 3.8411e-04
Epoch 68/100
7/7 _____ 0s 4ms/step - loss: 4.0888e-04
Epoch 69/100
7/7 _____ 0s 3ms/step - loss: 3.0997e-04
Epoch 70/100
7/7 _____ 0s 3ms/step - loss: 1.6669e-04
Epoch 71/100
7/7 _____ 0s 4ms/step - loss: 2.2517e-04
Epoch 72/100
7/7 _____ 0s 3ms/step - loss: 1.0623e-04
Epoch 73/100
7/7 _____ 0s 3ms/step - loss: 1.4488e-04
Epoch 74/100
7/7 _____ 0s 4ms/step - loss: 2.9294e-04
Epoch 75/100
7/7 _____ 0s 4ms/step - loss: 2.6550e-04
Epoch 76/100
7/7 _____ 0s 4ms/step - loss: 1.9894e-04
Epoch 77/100
7/7 _____ 0s 4ms/step - loss: 1.2361e-04
Epoch 78/100
7/7 _____ 0s 5ms/step - loss: 7.9886e-05
```

```
Epoch 79/100
7/7 _____ 0s 6ms/step - loss: 5.9101e-05
Epoch 80/100
7/7 _____ 0s 5ms/step - loss: 1.0901e-04
Epoch 81/100
7/7 _____ 0s 6ms/step - loss: 6.1116e-05
Epoch 82/100
7/7 _____ 0s 5ms/step - loss: 1.0688e-04
Epoch 83/100
7/7 _____ 0s 4ms/step - loss: 9.0334e-05
Epoch 84/100
7/7 _____ 0s 5ms/step - loss: 6.4809e-05
Epoch 85/100
7/7 _____ 0s 5ms/step - loss: 8.9256e-05
Epoch 86/100
7/7 _____ 0s 4ms/step - loss: 8.8408e-05
Epoch 87/100
7/7 _____ 0s 5ms/step - loss: 4.4358e-05
Epoch 88/100
7/7 _____ 0s 5ms/step - loss: 4.5070e-05
Epoch 89/100
7/7 _____ 0s 4ms/step - loss: 4.2587e-05
Epoch 90/100
7/7 _____ 0s 4ms/step - loss: 6.2594e-05
Epoch 91/100
7/7 _____ 0s 4ms/step - loss: 5.5651e-05
Epoch 92/100
7/7 _____ 0s 6ms/step - loss: 5.8087e-05
Epoch 93/100
7/7 _____ 0s 5ms/step - loss: 3.9013e-05
Epoch 94/100
7/7 _____ 0s 6ms/step - loss: 4.5478e-05
Epoch 95/100
7/7 _____ 0s 5ms/step - loss: 3.8964e-05
Epoch 96/100
7/7 _____ 0s 5ms/step - loss: 4.3546e-05
Epoch 97/100
7/7 _____ 0s 7ms/step - loss: 5.8468e-05
Epoch 98/100
7/7 _____ 0s 7ms/step - loss: 4.3040e-05
Epoch 99/100
7/7 _____ 0s 6ms/step - loss: 5.8045e-05
Epoch 100/100
7/7 _____ 0s 6ms/step - loss: 2.3155e-05
1/1 _____ 0s 286ms/step
Sequence: [ 1 2 4 8 16 32 64 128 256 512]
Predicted next number: 1004.71
```



```
import numpy as np
import tensorflow as tf

# Sample text data (first few lines of "Alice's Adventures in
Wonderland")
text = ("Alice was beginning to get very tired of sitting by her
sister on the bank, "
        "and of having nothing to do: once or twice she had peeped
into the book her sister "
        "was reading, but it had no pictures or conversations in it,
'and what is the use of "
        "a book,' thought Alice, 'without pictures or
conversations?'")

# Create a set of unique characters and a mapping from characters to
integers
chars = sorted(set(text))
char_to_int = {c: i for i, c in enumerate(chars)}
int_to_char = {i: c for i, c in enumerate(chars)}

# Prepare the input and output sequences
sequence_length = 5
X = []
y = []

for i in range(len(text) - sequence_length):
```

```

seq_in = text[i:i + sequence_length]
seq_out = text[i + sequence_length]
X.append([char_to_int[char] for char in seq_in])
y.append(char_to_int[seq_out])

X = np.array(X)
y = np.array(y)

# Reshape the input to be [samples, time steps, features]
X = X.reshape((X.shape[0], X.shape[1], 1))

# Normalize input
X = X / float(len(chars))

# Define the RNN model
model = tf.keras.Sequential([
    tf.keras.layers.SimpleRNN(100, input_shape=(X.shape[1], 1),
    return_sequences=False),
    tf.keras.layers.Dense(len(chars), activation='softmax')
])

/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/
rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

# Compile the model
model.compile(loss='sparse_categorical_crossentropy',
optimizer='adam', metrics=['accuracy'])
# Fit the model
model.fit(X, y, epochs=100, batch_size=32)

Epoch 1/100
10/10 ━━━━━━━━━━━ 2s 4ms/step - accuracy: 0.0841 - loss:
3.2235
Epoch 2/100
10/10 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.1979 - loss:
2.9969
Epoch 3/100
10/10 ━━━━━━━━━━━ 0s 4ms/step - accuracy: 0.1900 - loss:
2.9052
Epoch 4/100
10/10 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.1912 - loss:
2.8030
Epoch 5/100
10/10 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.2218 - loss:
2.7649
Epoch 6/100
10/10 ━━━━━━━━━━━ 0s 3ms/step - accuracy: 0.1860 - loss:

```

```
2.7317
Epoch 7/100
10/10 _____ 0s 3ms/step - accuracy: 0.2022 - loss:
2.7772
Epoch 8/100
10/10 _____ 0s 3ms/step - accuracy: 0.2007 - loss:
2.7453
Epoch 9/100
10/10 _____ 0s 3ms/step - accuracy: 0.1812 - loss:
2.7321
Epoch 10/100
10/10 _____ 0s 3ms/step - accuracy: 0.2121 - loss:
2.6904
Epoch 11/100
10/10 _____ 0s 3ms/step - accuracy: 0.2256 - loss:
2.6391
Epoch 12/100
10/10 _____ 0s 3ms/step - accuracy: 0.2449 - loss:
2.6409
Epoch 13/100
10/10 _____ 0s 3ms/step - accuracy: 0.1909 - loss:
2.7137
Epoch 14/100
10/10 _____ 0s 3ms/step - accuracy: 0.2002 - loss:
2.6657
Epoch 15/100
10/10 _____ 0s 3ms/step - accuracy: 0.1896 - loss:
2.6428
Epoch 16/100
10/10 _____ 0s 4ms/step - accuracy: 0.2075 - loss:
2.6467
Epoch 17/100
10/10 _____ 0s 4ms/step - accuracy: 0.2249 - loss:
2.6405
Epoch 18/100
10/10 _____ 0s 3ms/step - accuracy: 0.1802 - loss:
2.6873
Epoch 19/100
10/10 _____ 0s 3ms/step - accuracy: 0.2114 - loss:
2.6107
Epoch 20/100
10/10 _____ 0s 3ms/step - accuracy: 0.1793 - loss:
2.6806
Epoch 21/100
10/10 _____ 0s 3ms/step - accuracy: 0.2487 - loss:
2.5932
Epoch 22/100
10/10 _____ 0s 3ms/step - accuracy: 0.2025 - loss:
2.6357
```

```
Epoch 23/100
10/10 _____ 0s 3ms/step - accuracy: 0.2369 - loss:
2.5603
Epoch 24/100
10/10 _____ 0s 4ms/step - accuracy: 0.2092 - loss:
2.6333
Epoch 25/100
10/10 _____ 0s 3ms/step - accuracy: 0.2350 - loss:
2.6433
Epoch 26/100
10/10 _____ 0s 3ms/step - accuracy: 0.2341 - loss:
2.5453
Epoch 27/100
10/10 _____ 0s 3ms/step - accuracy: 0.2221 - loss:
2.5735
Epoch 28/100
10/10 _____ 0s 3ms/step - accuracy: 0.2496 - loss:
2.5225
Epoch 29/100
10/10 _____ 0s 4ms/step - accuracy: 0.2560 - loss:
2.5658
Epoch 30/100
10/10 _____ 0s 3ms/step - accuracy: 0.2385 - loss:
2.5390
Epoch 31/100
10/10 _____ 0s 3ms/step - accuracy: 0.2085 - loss:
2.5650
Epoch 32/100
10/10 _____ 0s 3ms/step - accuracy: 0.2350 - loss:
2.5842
Epoch 33/100
10/10 _____ 0s 3ms/step - accuracy: 0.1927 - loss:
2.5346
Epoch 34/100
10/10 _____ 0s 3ms/step - accuracy: 0.2275 - loss:
2.5615
Epoch 35/100
10/10 _____ 0s 3ms/step - accuracy: 0.2549 - loss:
2.5191
Epoch 36/100
10/10 _____ 0s 3ms/step - accuracy: 0.2185 - loss:
2.6064
Epoch 37/100
10/10 _____ 0s 3ms/step - accuracy: 0.2552 - loss:
2.4991
Epoch 38/100
10/10 _____ 0s 3ms/step - accuracy: 0.2245 - loss:
2.5382
Epoch 39/100
```



```
10/10 _____ 0s 3ms/step - accuracy: 0.2098 - loss:
2.5640
Epoch 40/100
10/10 _____ 0s 3ms/step - accuracy: 0.2141 - loss:
2.5585
Epoch 41/100
10/10 _____ 0s 4ms/step - accuracy: 0.2023 - loss:
2.5602
Epoch 42/100
10/10 _____ 0s 8ms/step - accuracy: 0.2491 - loss:
2.4850
Epoch 43/100
10/10 _____ 0s 5ms/step - accuracy: 0.2801 - loss:
2.4796
Epoch 44/100
10/10 _____ 0s 5ms/step - accuracy: 0.2717 - loss:
2.4524
Epoch 45/100
10/10 _____ 0s 7ms/step - accuracy: 0.3032 - loss:
2.3967
Epoch 46/100
10/10 _____ 0s 5ms/step - accuracy: 0.2677 - loss:
2.4582
Epoch 47/100
10/10 _____ 0s 5ms/step - accuracy: 0.2843 - loss:
2.4398
Epoch 48/100
10/10 _____ 0s 5ms/step - accuracy: 0.2475 - loss:
2.5200
Epoch 49/100
10/10 _____ 0s 5ms/step - accuracy: 0.2938 - loss:
2.4003
Epoch 50/100
10/10 _____ 0s 5ms/step - accuracy: 0.2442 - loss:
2.3965
Epoch 51/100
10/10 _____ 0s 4ms/step - accuracy: 0.2323 - loss:
2.4644
Epoch 52/100
10/10 _____ 0s 5ms/step - accuracy: 0.2588 - loss:
2.4267
Epoch 53/100
10/10 _____ 0s 5ms/step - accuracy: 0.2728 - loss:
2.4201
Epoch 54/100
10/10 _____ 0s 5ms/step - accuracy: 0.2746 - loss:
2.4166
Epoch 55/100
10/10 _____ 0s 6ms/step - accuracy: 0.2270 - loss:
```

```
2.4556
Epoch 56/100
10/10 _____ 0s 5ms/step - accuracy: 0.2566 - loss:
2.3732
Epoch 57/100
10/10 _____ 0s 5ms/step - accuracy: 0.2262 - loss:
2.4484
Epoch 58/100
10/10 _____ 0s 5ms/step - accuracy: 0.3107 - loss:
2.3380
Epoch 59/100
10/10 _____ 0s 5ms/step - accuracy: 0.2877 - loss:
2.3551
Epoch 60/100
10/10 _____ 0s 5ms/step - accuracy: 0.3131 - loss:
2.3117
Epoch 61/100
10/10 _____ 0s 7ms/step - accuracy: 0.2247 - loss:
2.4157
Epoch 62/100
10/10 _____ 0s 6ms/step - accuracy: 0.3001 - loss:
2.2673
Epoch 63/100
10/10 _____ 0s 8ms/step - accuracy: 0.3228 - loss:
2.2488
Epoch 64/100
10/10 _____ 0s 4ms/step - accuracy: 0.3110 - loss:
2.3056
Epoch 65/100
10/10 _____ 0s 4ms/step - accuracy: 0.2755 - loss:
2.3308
Epoch 66/100
10/10 _____ 0s 3ms/step - accuracy: 0.2775 - loss:
2.3136
Epoch 67/100
10/10 _____ 0s 3ms/step - accuracy: 0.3025 - loss:
2.2642
Epoch 68/100
10/10 _____ 0s 3ms/step - accuracy: 0.3318 - loss:
2.2911
Epoch 69/100
10/10 _____ 0s 4ms/step - accuracy: 0.2705 - loss:
2.3449
Epoch 70/100
10/10 _____ 0s 3ms/step - accuracy: 0.2813 - loss:
2.2902
Epoch 71/100
10/10 _____ 0s 3ms/step - accuracy: 0.3046 - loss:
2.2783
```

```
Epoch 72/100
10/10 _____ 0s 4ms/step - accuracy: 0.3004 - loss:
2.2713
Epoch 73/100
10/10 _____ 0s 3ms/step - accuracy: 0.3237 - loss:
2.2536
Epoch 74/100
10/10 _____ 0s 3ms/step - accuracy: 0.2962 - loss:
2.2556
Epoch 75/100
10/10 _____ 0s 3ms/step - accuracy: 0.3389 - loss:
2.1917
Epoch 76/100
10/10 _____ 0s 4ms/step - accuracy: 0.3212 - loss:
2.2385
Epoch 77/100
10/10 _____ 0s 4ms/step - accuracy: 0.2884 - loss:
2.2298
Epoch 78/100
10/10 _____ 0s 4ms/step - accuracy: 0.2983 - loss:
2.2432
Epoch 79/100
10/10 _____ 0s 5ms/step - accuracy: 0.3084 - loss:
2.1828
Epoch 80/100
10/10 _____ 0s 4ms/step - accuracy: 0.3265 - loss:
2.2021
Epoch 81/100
10/10 _____ 0s 3ms/step - accuracy: 0.2925 - loss:
2.2351
Epoch 82/100
10/10 _____ 0s 3ms/step - accuracy: 0.2950 - loss:
2.2043
Epoch 83/100
10/10 _____ 0s 3ms/step - accuracy: 0.3387 - loss:
2.1649
Epoch 84/100
10/10 _____ 0s 3ms/step - accuracy: 0.3124 - loss:
2.2189
Epoch 85/100
10/10 _____ 0s 3ms/step - accuracy: 0.2900 - loss:
2.1531
Epoch 86/100
10/10 _____ 0s 3ms/step - accuracy: 0.2761 - loss:
2.2035
Epoch 87/100
10/10 _____ 0s 3ms/step - accuracy: 0.3335 - loss:
2.1322
Epoch 88/100
```

```

10/10 _____ 0s 3ms/step - accuracy: 0.2917 - loss:
2.1624
Epoch 89/100
10/10 _____ 0s 3ms/step - accuracy: 0.3396 - loss:
2.0920
Epoch 90/100
10/10 _____ 0s 4ms/step - accuracy: 0.3217 - loss:
2.1489
Epoch 91/100
10/10 _____ 0s 4ms/step - accuracy: 0.3135 - loss:
2.1676
Epoch 92/100
10/10 _____ 0s 3ms/step - accuracy: 0.3420 - loss:
2.1100
Epoch 93/100
10/10 _____ 0s 3ms/step - accuracy: 0.3350 - loss:
2.1146
Epoch 94/100
10/10 _____ 0s 3ms/step - accuracy: 0.3460 - loss:
2.0988
Epoch 95/100
10/10 _____ 0s 3ms/step - accuracy: 0.3963 - loss:
1.9904
Epoch 96/100
10/10 _____ 0s 3ms/step - accuracy: 0.3570 - loss:
2.0936
Epoch 97/100
10/10 _____ 0s 4ms/step - accuracy: 0.3925 - loss:
2.0632
Epoch 98/100
10/10 _____ 0s 3ms/step - accuracy: 0.3687 - loss:
2.0485
Epoch 99/100
10/10 _____ 0s 3ms/step - accuracy: 0.3577 - loss:
2.0233
Epoch 100/100
10/10 _____ 0s 3ms/step - accuracy: 0.3363 - loss:
2.0386

```

```
<keras.src.callbacks.history.History at 0x7bfc16ad59c0>
```

```
# Function to predict the next character
```

```

def predict_next_char(model, input_seq):
    input_data = np.array([char_to_int[char] for char in input_seq])
    input_data = input_data.reshape((1, len(input_seq), 1)) /
float(len(chars))

    prediction = model.predict(input_data, verbose=0)
    index = np.argmax(prediction)

```

```
    return int_to_char[index]
```

```
# Example prediction
```

```
input_sequence = "the book"
```

```
next_char = predict_next_char(model, input_sequence)
```

```
print(f"The next character after '{input_sequence}' is '{next_char}'")
```

```
The next character after 'the book' is 'p'
```

Assignment-9

Q1. import numpy **as** np

import pandas **as** pd

import matplotlib.pyplot **as** plt

df = pd.read_csv('LSTM.csv', index_col='Date', parse_dates=**True**)

df.index.freq = 'MS'

Q2. Find head()

Q3. Plot prediction Vs Year.

Q4. Select first 156 rows for training and 156 onwards for test.

Q5. Normalize the data.

Q6. Add following code:

from keras.preprocessing.sequence **import** TimeseriesGenerator

define generator model

n_inputs = 3

n_features = 1

generator = TimeseriesGenerator(scaled_train, scaled_train, length=n_inputs,
batch_size=1)

Q7. Find output:

Production	
Date	
1975-01-01	834
1975-02-01	782
1975-03-01	892
1975-04-01	903
1975-05-01	966

```
X,y = generator[1]
```

```
print(f'Given the array: \n {X.flatten()}')
```

```
print(f'Predict this y: \n {y}')
```

Q8. do the same thing but with 12 months of data

```
n_inputs = 12
```

```
generator = TimeseriesGenerator(scaled_train, scaled_train, length=n_inputs,  
    batch_size=1)
```

Q9. Use LSTM with 100 inputs with ReLU and use adam and loss=mse.

Q10. Find model summary

Q11. fit model with epocs=50.

Q12. Plot loss curve.

Q13. Find the following output


Q14. Plot production vs prediction.

Q15. calculate the error of the model.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\Asus\Downloads\LSTM.csv",
index_col='Date', parse_dates=True)
df.index.freq = 'MS'

df.head()

# Plot the production data vs. year
plt.figure(figsize=(10, 5))
plt.plot(df.index, df['Production'], label="Production")
plt.xlabel("Year")
plt.ylabel("Production Value")
plt.title("Production vs Year")
plt.legend()
plt.show()
```




```

train = df.iloc[:156]
test = df.iloc[156:]
print(train.size)
print(test.size)

156
12

# Q6: Normalize the data
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaled_train = scaler.fit_transform(train)
scaled_test = scaler.transform(test)

# Q7-8: Define Timeseries Generator for both 3-month and 12-month
inputs
from sklearn.preprocessing import MinMaxScaler
from keras.models import Sequential
from keras.layers import LSTM, Dense
from tensorflow.keras.preprocessing.sequence import
TimeseriesGenerator
from sklearn.metrics import mean_squared_error
# 3-month generator
n_inputs_3 = 3
n_inputs_12 = 12
generator_3 = TimeseriesGenerator(scaled_train, scaled_train,
length=n_inputs_3, batch_size=1)
generator_12 = TimeseriesGenerator(scaled_train, scaled_train,
length=n_inputs_12, batch_size=1)
# Sample output for 3-month generator
X_3, y_3 = generator_3[1]
print(f'Given the array (3 months): \n {X_3.flatten()}')
print(f'Predict this y: \n {y_3}')
# Sample output for 12-month generator
X, y = generator_12[1]
print(f'Given the array (12 months): \n {X.flatten()}')
print(f'Predict this y: \n {y}')

Given the array (3 months):
[0.01923077 0.20913462 0.24759615]
Predict this y:
[[0.41826923]]
Given the array (12 months):
[0.01923077 0.20913462 0.24759615 0.41826923 0.34615385 0.20913462
0.11057692 0.03605769 0.05769231 0.          0.06971154 0.11298077]
Predict this y:
[[0.03125]]

# Q9-10: Define the LSTM model
model = Sequential([

```

```

    LSTM(100, activation='relu', input_shape=(n_inputs_12, 1)), #
    input_shape for 12-month generator
    Dense(1)
])
model.compile(optimizer='adam', loss='mse')
# Display model summary
print("Model Summary:")
model.summary()

```

Model Summary:

C:\Users\Asus\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\layers\rnn\rnn.py:204: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

    super().__init__(**kwargs)

```

Model: "sequential_1"

Layer (type) Param #	Output Shape
lstm_1 (LSTM) 40,800	(None, 100)
dense_1 (Dense) 101	(None, 1)

Total params: 40,901 (159.77 KB)

Trainable params: 40,901 (159.77 KB)

Non-trainable params: 0 (0.00 B)

Q11: Fit the model with 50 epochs

```

generator_12 = TimeseriesGenerator(scaled_train, scaled_train,
length=n_inputs_12, batch_size=1)
history = model.fit(generator_12, epochs=50, verbose=1)

```

C:\Users\Asus\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not

```
pass these arguments to `fit()`, as they will be ignored.  
self._warn_if_super_not_called()
```

```
Epoch 1/50  
144/144 _____ 14s 15ms/step - loss: 0.0983  
Epoch 2/50  
144/144 _____ 2s 15ms/step - loss: 0.0178  
Epoch 3/50  
144/144 _____ 3s 19ms/step - loss: 0.0127  
Epoch 4/50  
144/144 _____ 3s 17ms/step - loss: 0.0146  
Epoch 5/50  
144/144 _____ 3s 21ms/step - loss: 0.0084  
Epoch 6/50  
144/144 _____ 3s 18ms/step - loss: 0.0066  
Epoch 7/50  
144/144 _____ 3s 18ms/step - loss: 0.0052  
Epoch 8/50  
144/144 _____ 3s 18ms/step - loss: 0.0087  
Epoch 9/50  
144/144 _____ 2s 16ms/step - loss: 0.0041  
Epoch 10/50  
144/144 _____ 2s 15ms/step - loss: 0.0049  
Epoch 11/50  
144/144 _____ 2s 15ms/step - loss: 0.0046  
Epoch 12/50  
144/144 _____ 4s 23ms/step - loss: 0.0045  
Epoch 13/50  
144/144 _____ 3s 22ms/step - loss: 0.0048  
Epoch 14/50  
144/144 _____ 5s 18ms/step - loss: 0.0042  
Epoch 15/50  
144/144 _____ 6s 21ms/step - loss: 0.0028  
Epoch 16/50  
144/144 _____ 3s 17ms/step - loss: 0.0033  
Epoch 17/50  
144/144 _____ 3s 19ms/step - loss: 0.0048  
Epoch 18/50  
144/144 _____ 3s 21ms/step - loss: 0.0031  
Epoch 19/50  
144/144 _____ 3s 21ms/step - loss: 0.0031  
Epoch 20/50  
144/144 _____ 3s 16ms/step - loss: 0.0036  
Epoch 21/50  
144/144 _____ 3s 17ms/step - loss: 0.0021  
Epoch 22/50  
144/144 _____ 3s 18ms/step - loss: 0.0054  
Epoch 23/50  
144/144 _____ 2s 16ms/step - loss: 0.0037  
Epoch 24/50
```

144/144	_____	2s 15ms/step	- loss: 0.0039
Epoch 25/50			
144/144	_____	2s 15ms/step	- loss: 0.0031
Epoch 26/50			
144/144	_____	2s 16ms/step	- loss: 0.0028
Epoch 27/50			
144/144	_____	2s 14ms/step	- loss: 0.0027
Epoch 28/50			
144/144	_____	3s 18ms/step	- loss: 0.0043
Epoch 29/50			
144/144	_____	6s 20ms/step	- loss: 0.0023
Epoch 30/50			
144/144	_____	3s 19ms/step	- loss: 0.0045
Epoch 31/50			
144/144	_____	2s 16ms/step	- loss: 0.0032
Epoch 32/50			
144/144	_____	2s 15ms/step	- loss: 0.0019
Epoch 33/50			
144/144	_____	2s 14ms/step	- loss: 0.0023
Epoch 34/50			
144/144	_____	3s 16ms/step	- loss: 0.0028
Epoch 35/50			
144/144	_____	2s 15ms/step	- loss: 0.0044
Epoch 36/50			
144/144	_____	2s 15ms/step	- loss: 0.0027
Epoch 37/50			
144/144	_____	3s 16ms/step	- loss: 0.0024
Epoch 38/50			
144/144	_____	2s 15ms/step	- loss: 0.0024
Epoch 39/50			
144/144	_____	2s 15ms/step	- loss: 0.0020
Epoch 40/50			
144/144	_____	3s 15ms/step	- loss: 0.0030
Epoch 41/50			
144/144	_____	2s 15ms/step	- loss: 0.0020
Epoch 42/50			
144/144	_____	2s 15ms/step	- loss: 0.0026
Epoch 43/50			
144/144	_____	2s 15ms/step	- loss: 0.0020
Epoch 44/50			
144/144	_____	2s 15ms/step	- loss: 0.0022
Epoch 45/50			
144/144	_____	2s 15ms/step	- loss: 0.0025
Epoch 46/50			
144/144	_____	2s 15ms/step	- loss: 0.0022
Epoch 47/50			
144/144	_____	2s 15ms/step	- loss: 0.0022
Epoch 48/50			
144/144	_____	2s 15ms/step	- loss: 0.0019

```
Epoch 49/50
144/144 ————— 3s 16ms/step - loss: 0.0022
Epoch 50/50
144/144 ————— 3s 17ms/step - loss: 0.0027
```

```
# Q12: Plot the training loss curve
plt.plot(history.history['loss'])
plt.title("Training Loss Curve")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.show()


# Q13-14: Make predictions
predictions = []
batch = scaled_train[-n_inputs_12:] # Last 12 months from training as
initial batch
current_batch = batch.reshape((1, n_inputs_12, 1))

for i in range(len(test)):
    pred = model.predict(current_batch)[0]
    predictions.append(pred)
    current_batch = np.append(current_batch[:, 1:, :], [[pred]],
axis=1)


# Inverse transform predictions
predictions = scaler.inverse_transform(predictions)

# Plot production vs prediction
plt.figure(figsize=(10, 5))
plt.plot(df.index[:156], train, label="Training Data")
plt.plot(df.index[156:], test, label="Actual Data")
plt.plot(df.index[156:], predictions, label="Predictions")
plt.xlabel("Year")
plt.ylabel("Production Value")
plt.title("Production vs Prediction")
plt.legend()
plt.show()

# Q15: Calculate and print Mean Squared Error (MSE)
mse = mean_squared_error(test, predictions)
print("Model Error (MSE):", mse)
```



1/1	1s 834ms/step
1/1	0s 129ms/step
1/1	0s 109ms/step
1/1	0s 113ms/step
1/1	0s 99ms/step
1/1	0s 133ms/step
1/1	0s 112ms/step
1/1	0s 144ms/step
1/1	0s 131ms/step
1/1	0s 102ms/step
1/1	0s 127ms/step
1/1	0s 94ms/step



Model Error (MSE): 969.8440586116493